



Kubernetes

Part 1



Cloud Native Application

- Cloud native computing
 - Segment application into microservices
 - Package each part of into its own container
 - Dynamically orchestrate those containers to optimize resource utilization



CLOUD NATIVE
COMPUTING FOUNDATION

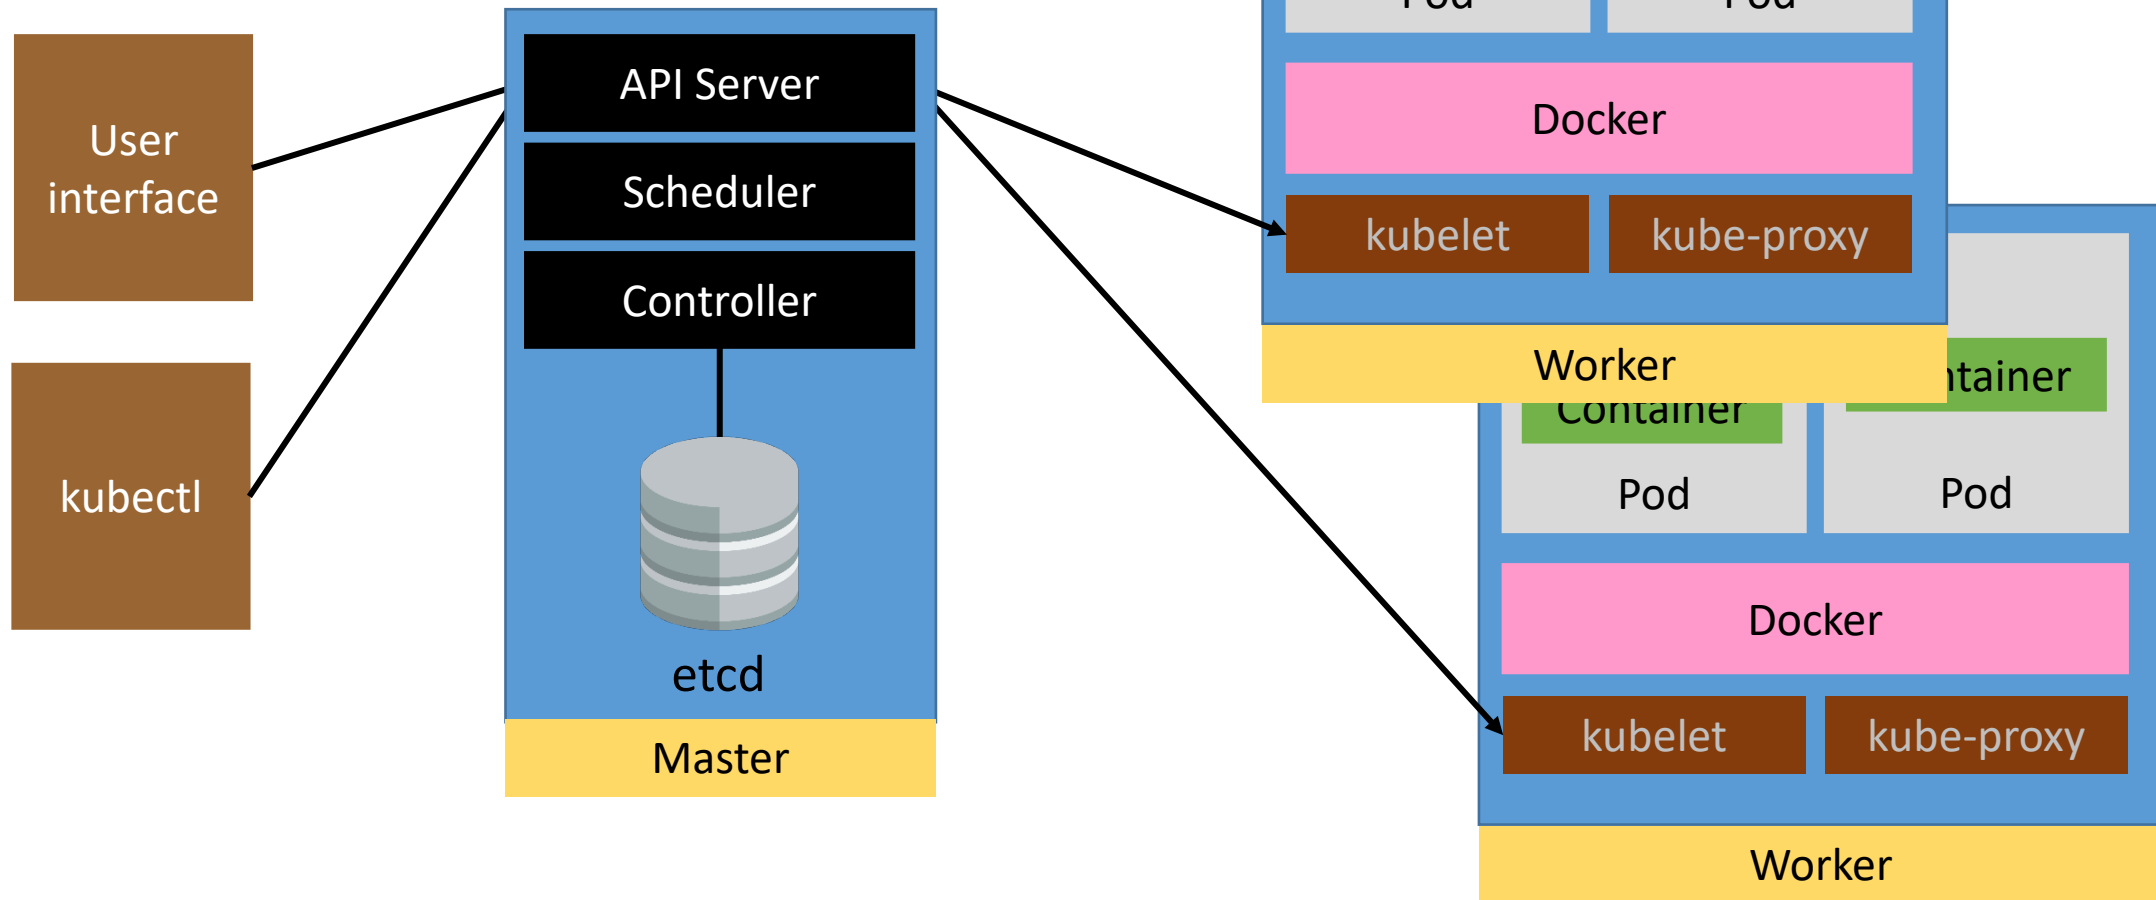


What is Kubernetes?

- Container orchestrator
 - Schedules and deploys containers
 - Recover from failure, keeping the actual state and desired state of an application in sync
 - Provides basic monitoring, logging, health checks
 - Enables containers to talk to each other
 - Scale workloads
- Project that started off as an internal Google project for managing containers
- Provides the same API across all cloud providers
 - Free from the underlying cloud platform



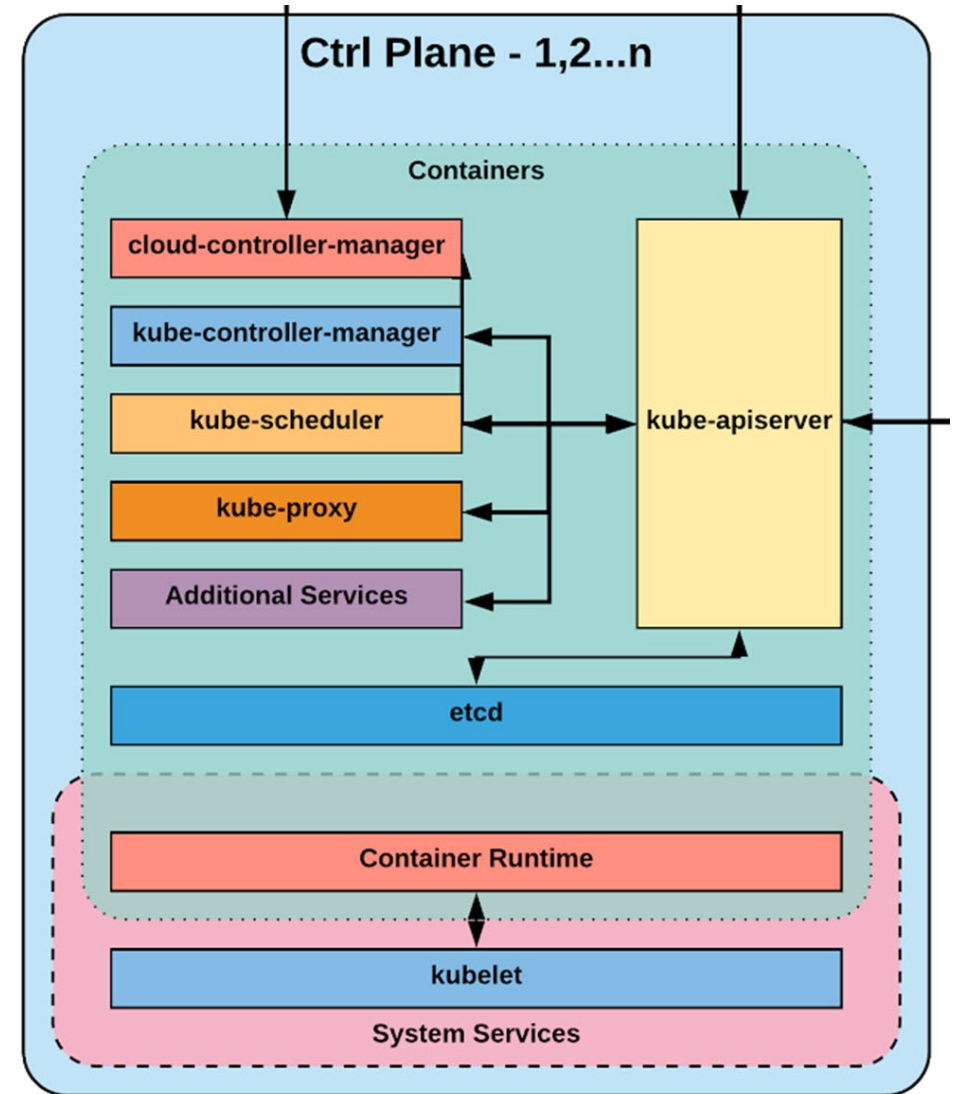
Architecture






Control Plane Components

- API Server – receives REST request to create services, deploy Pods, etc
- Controller manager – runs a number of process to manages the cluster. Monitors the state of the cluster and steers the cluster towards the desired state
- Scheduler – evaluates workload requirements and schedule it on matching resources
- Etcd – a distributed key/value storage; used to store the cluster's state





Node Components

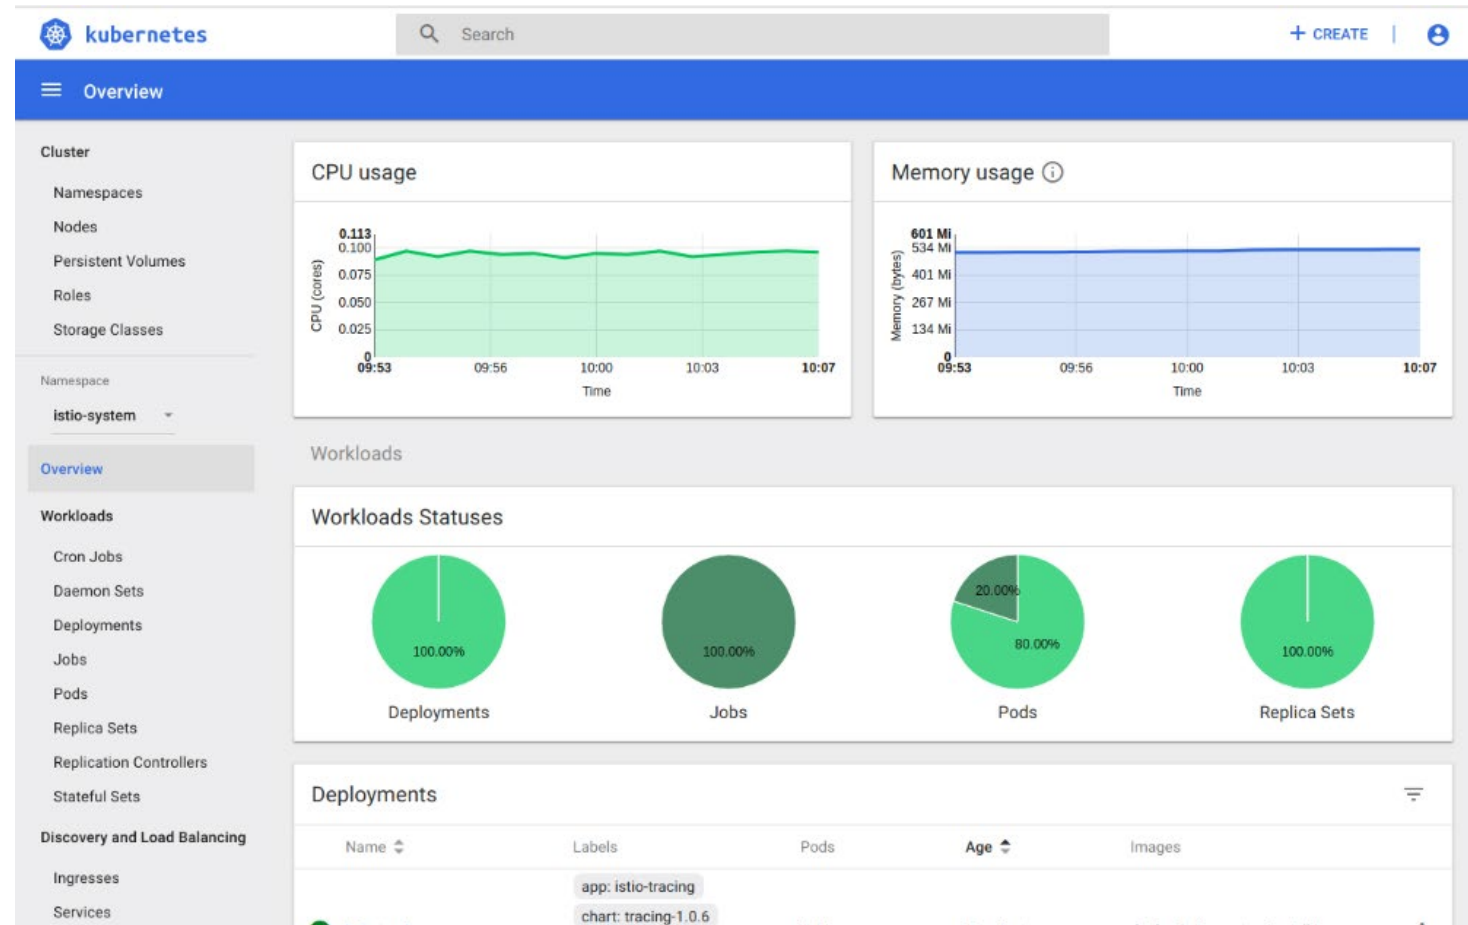


- Kubelet – communicates with the API server in master. Takes orders from masters to schedule and manage Pods. Also reports the health of the Pods to the master node
- Kube-proxy – handles' the Pods networking. Performs connection forwarding and load balancing for services
- Container runtime – use to manage containers. Docker in our case



Interacting with Kubernetes

- Native Kubernetes dashboard
 - Not dashboard provided by cloud providers
 - May need to install
- **kubectl**
 - Command line
- Programmatic
 - REST API
 - Language specific client libraries



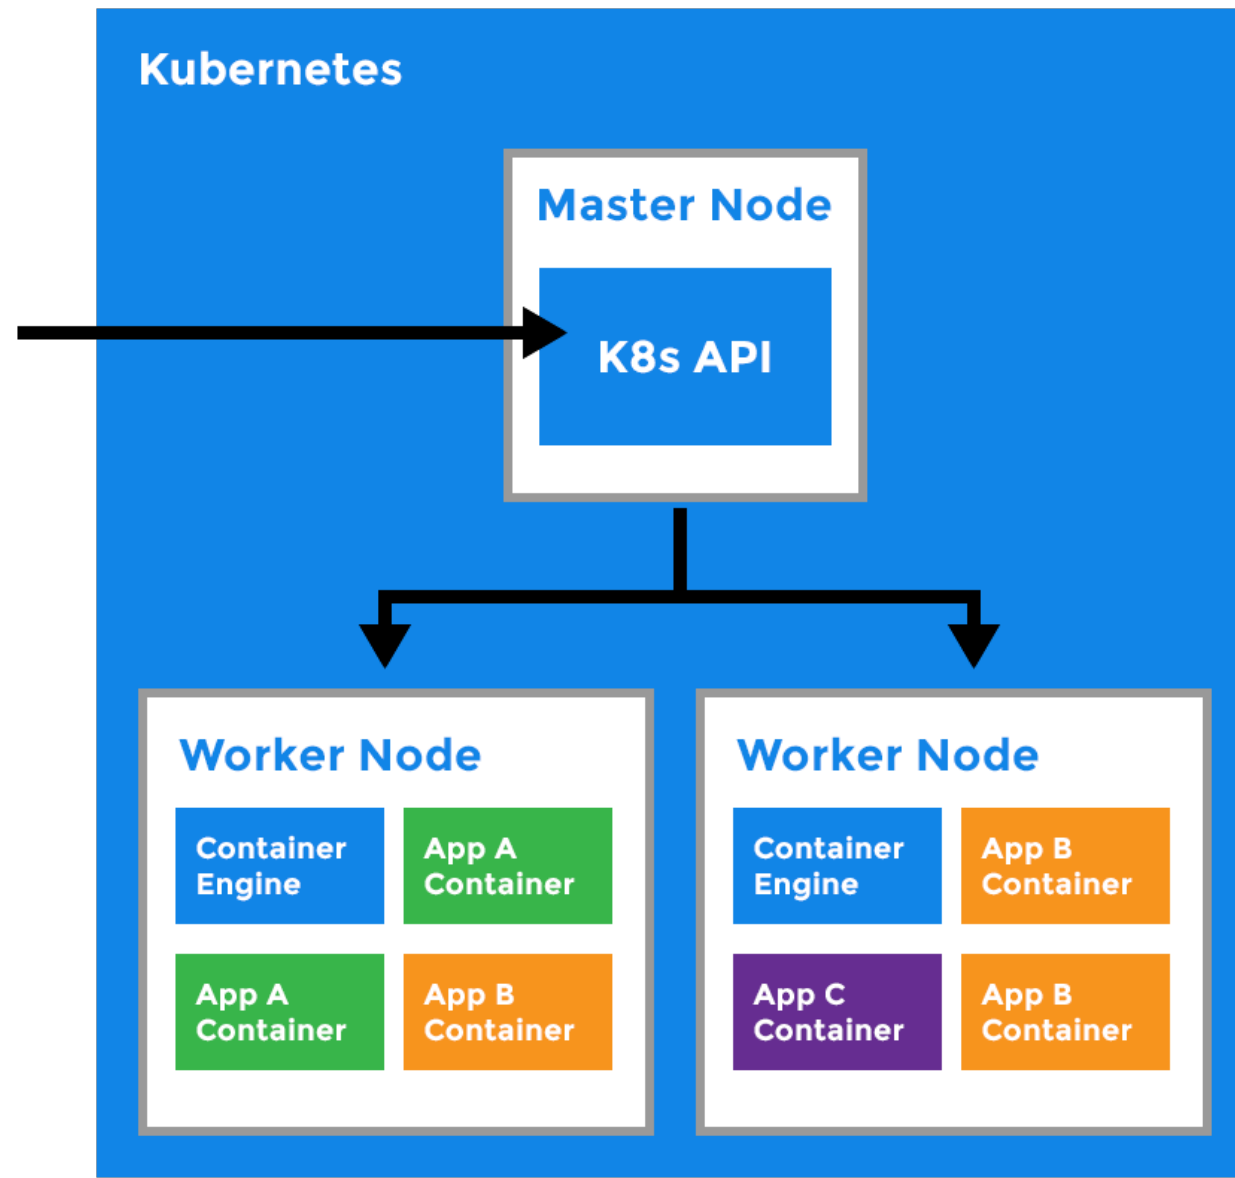


kubectl

- Command line tool for interacting with the cluster
- Gets cluster information from configuration file
 - Default location
`$HOME/.kube/config`
- If file is in different location
 - Use `--kubeconfig` option
 - Set `KUBECONFIG` environment variable



kubectl





Kubernetes Command

- Creating a resource

```
kubectl apply -f <yaml_file>
```

- Getting a resource

```
kubectl get <resource_type> <resource_name>
```

- Detailed information

```
kubectl describe <resource_type> <resource_name>
```

- Delete a resource

```
kubectl delete <resource_type> <resource_name>
```

```
kubectl delete -f <yaml_file>
```

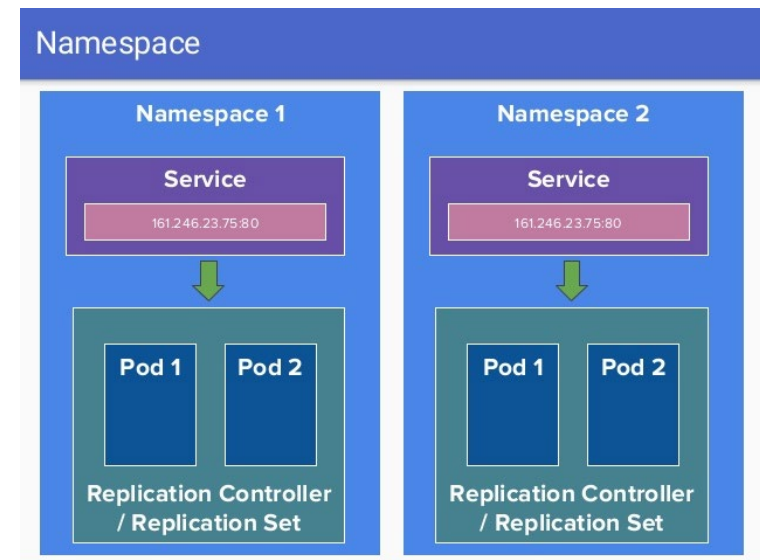
Resource type

pod	Pod
deploy	Deployment
svc	Service
ing	Ingress
pvc	Persistent volume claim



Namespace

- A single Kubernetes cluster should be able to run application from multiple users
 - For security reasons
 - Each user/application should only be allowed to access their own resources
 - Different policies to regulate their access
 - Kubernetes uses namespace to isolate users and applications
 - Can restrict resources, access inside a namespace
- Default namespaces
 - `default` - the namespace to that Kubernetes put your deployment into if you did not specify any namespace
 - `kube-system` - for system
 - `kube-public` - accessible by all user





Defining and Using Namespace

- Create the namespace

```
kubectl apply -f namespace.yml
```

- Specify the namespace with -n

```
kubectl -n <name> <command>
```

- Delete namespace

```
kubectl delete namespace/<name>
```

```
apiVersion: v1  
kind: Namespace  
metadata:
```


```
  name: myapp
```

Namespace

A black arrow pointing from the word "Namespace" to the text "myapp" in the previous block.




Pod



- Is a “unit of work” in Kubernetes
 - Smallest schedulable entity
- Contains one or more containers
- Containers in a Pod share the same namespace
 - Allow containers inside a Pod to communicate with each other via localhost
- Containers in a Pod are either all running or they are not
 - Can never have a Pod with one or more failed containers
- Each Pod has an IP address
- Pods are ephemeral
 - They can die
 - They can be rescheduled by Kubernetes to another node



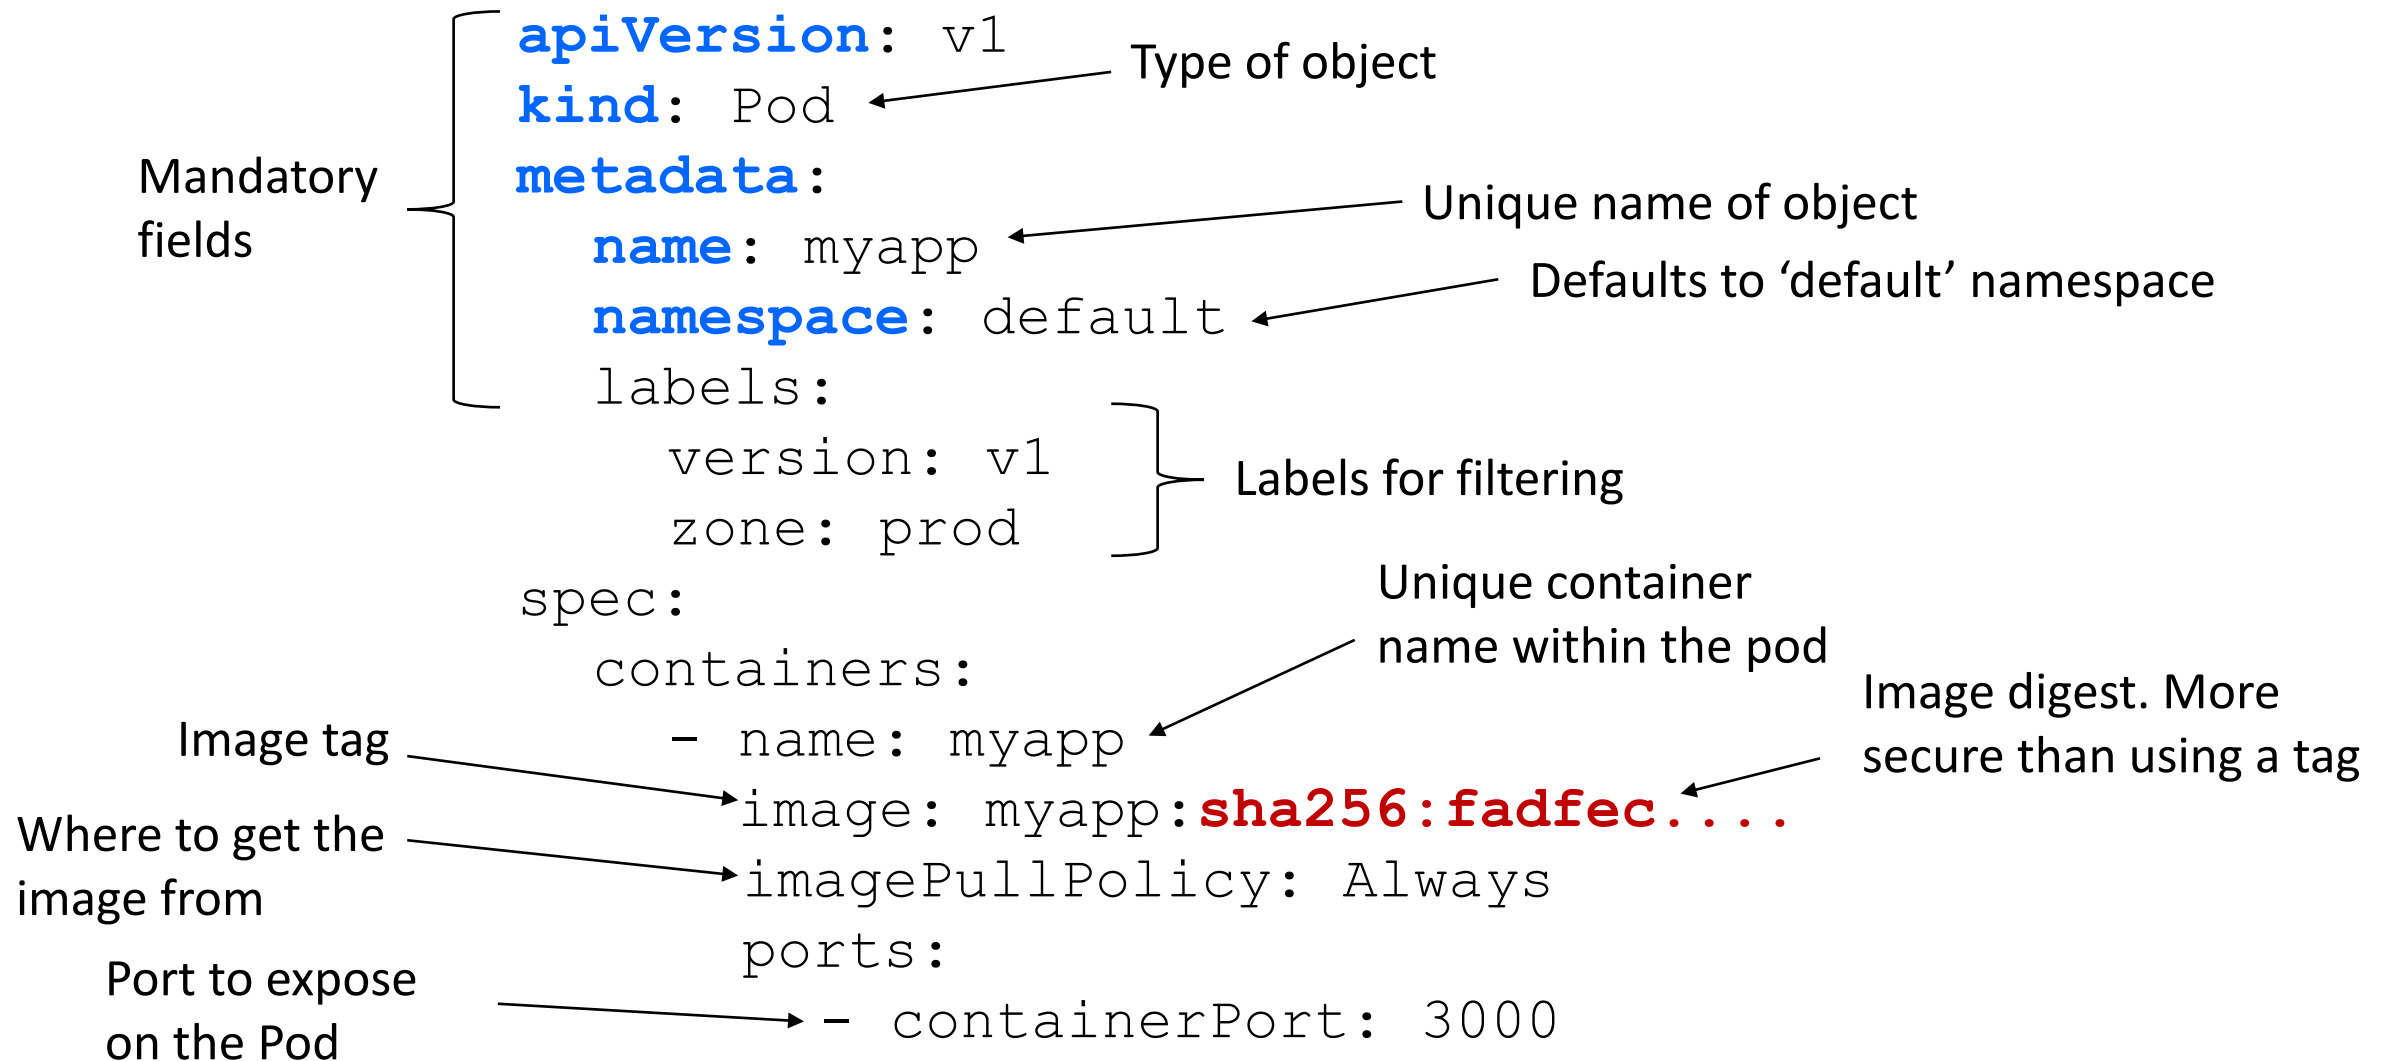
Example of Multi Container Pod



- Two containers
 - Web server displays content from a volume
 - File puller to ensure that the volume has the latest content by syncing it with some remote master
- The two containers are tightly coupled and should be scheduled as a Pod



Defining a Pod





Pod Management

- **Create a Pod**

```
kubectl apply -f pod.yml
```

- **View all Pods**

```
kubectl get pods -o wide
```

```
kubectl get pods -o yaml
```

- **Detail information about a pod**

```
kubectl describe pods myapp-pod
```

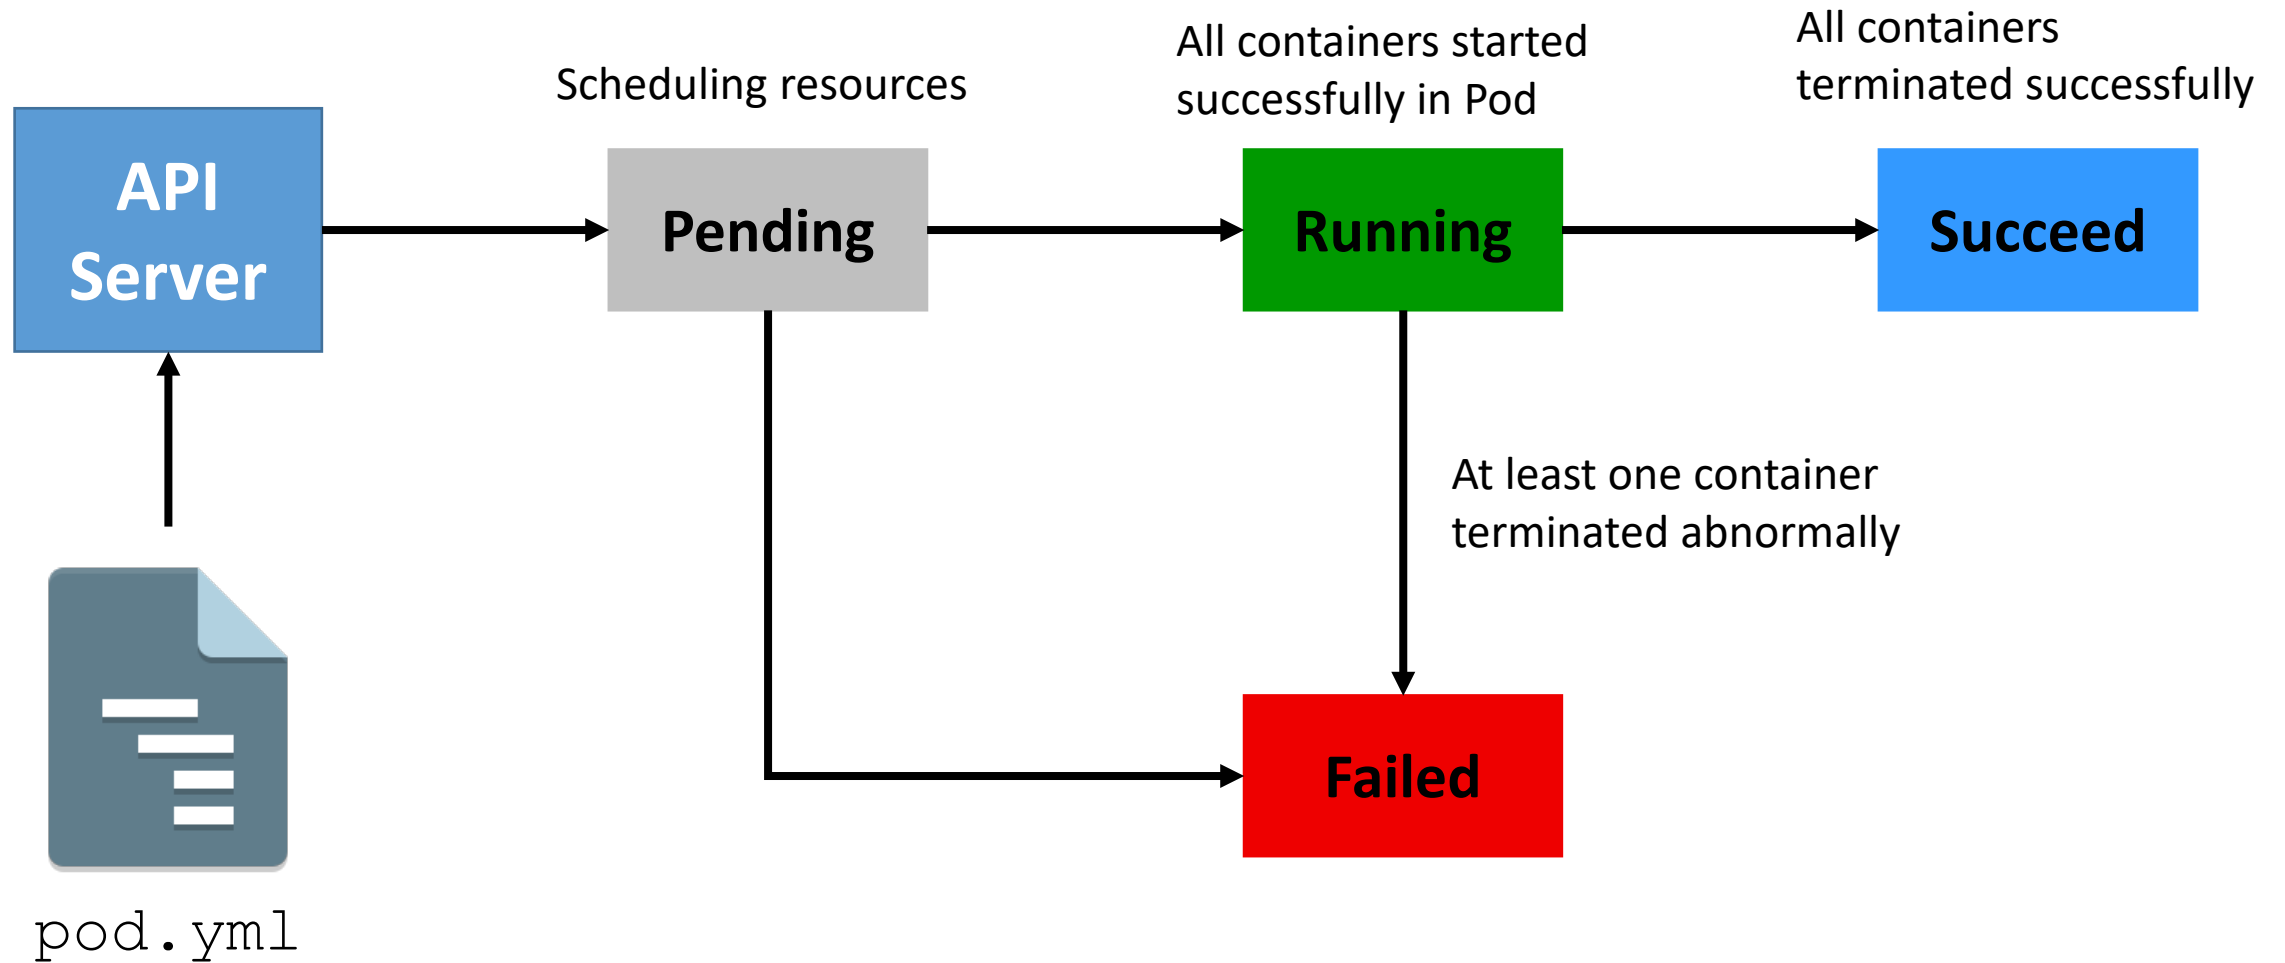
- **Delete a Pod**

```
kubectl delete -f pod.yml
```

```
kubectl delete pod myapp-pod
```



Pod Lifecycle





Accessing the Pod

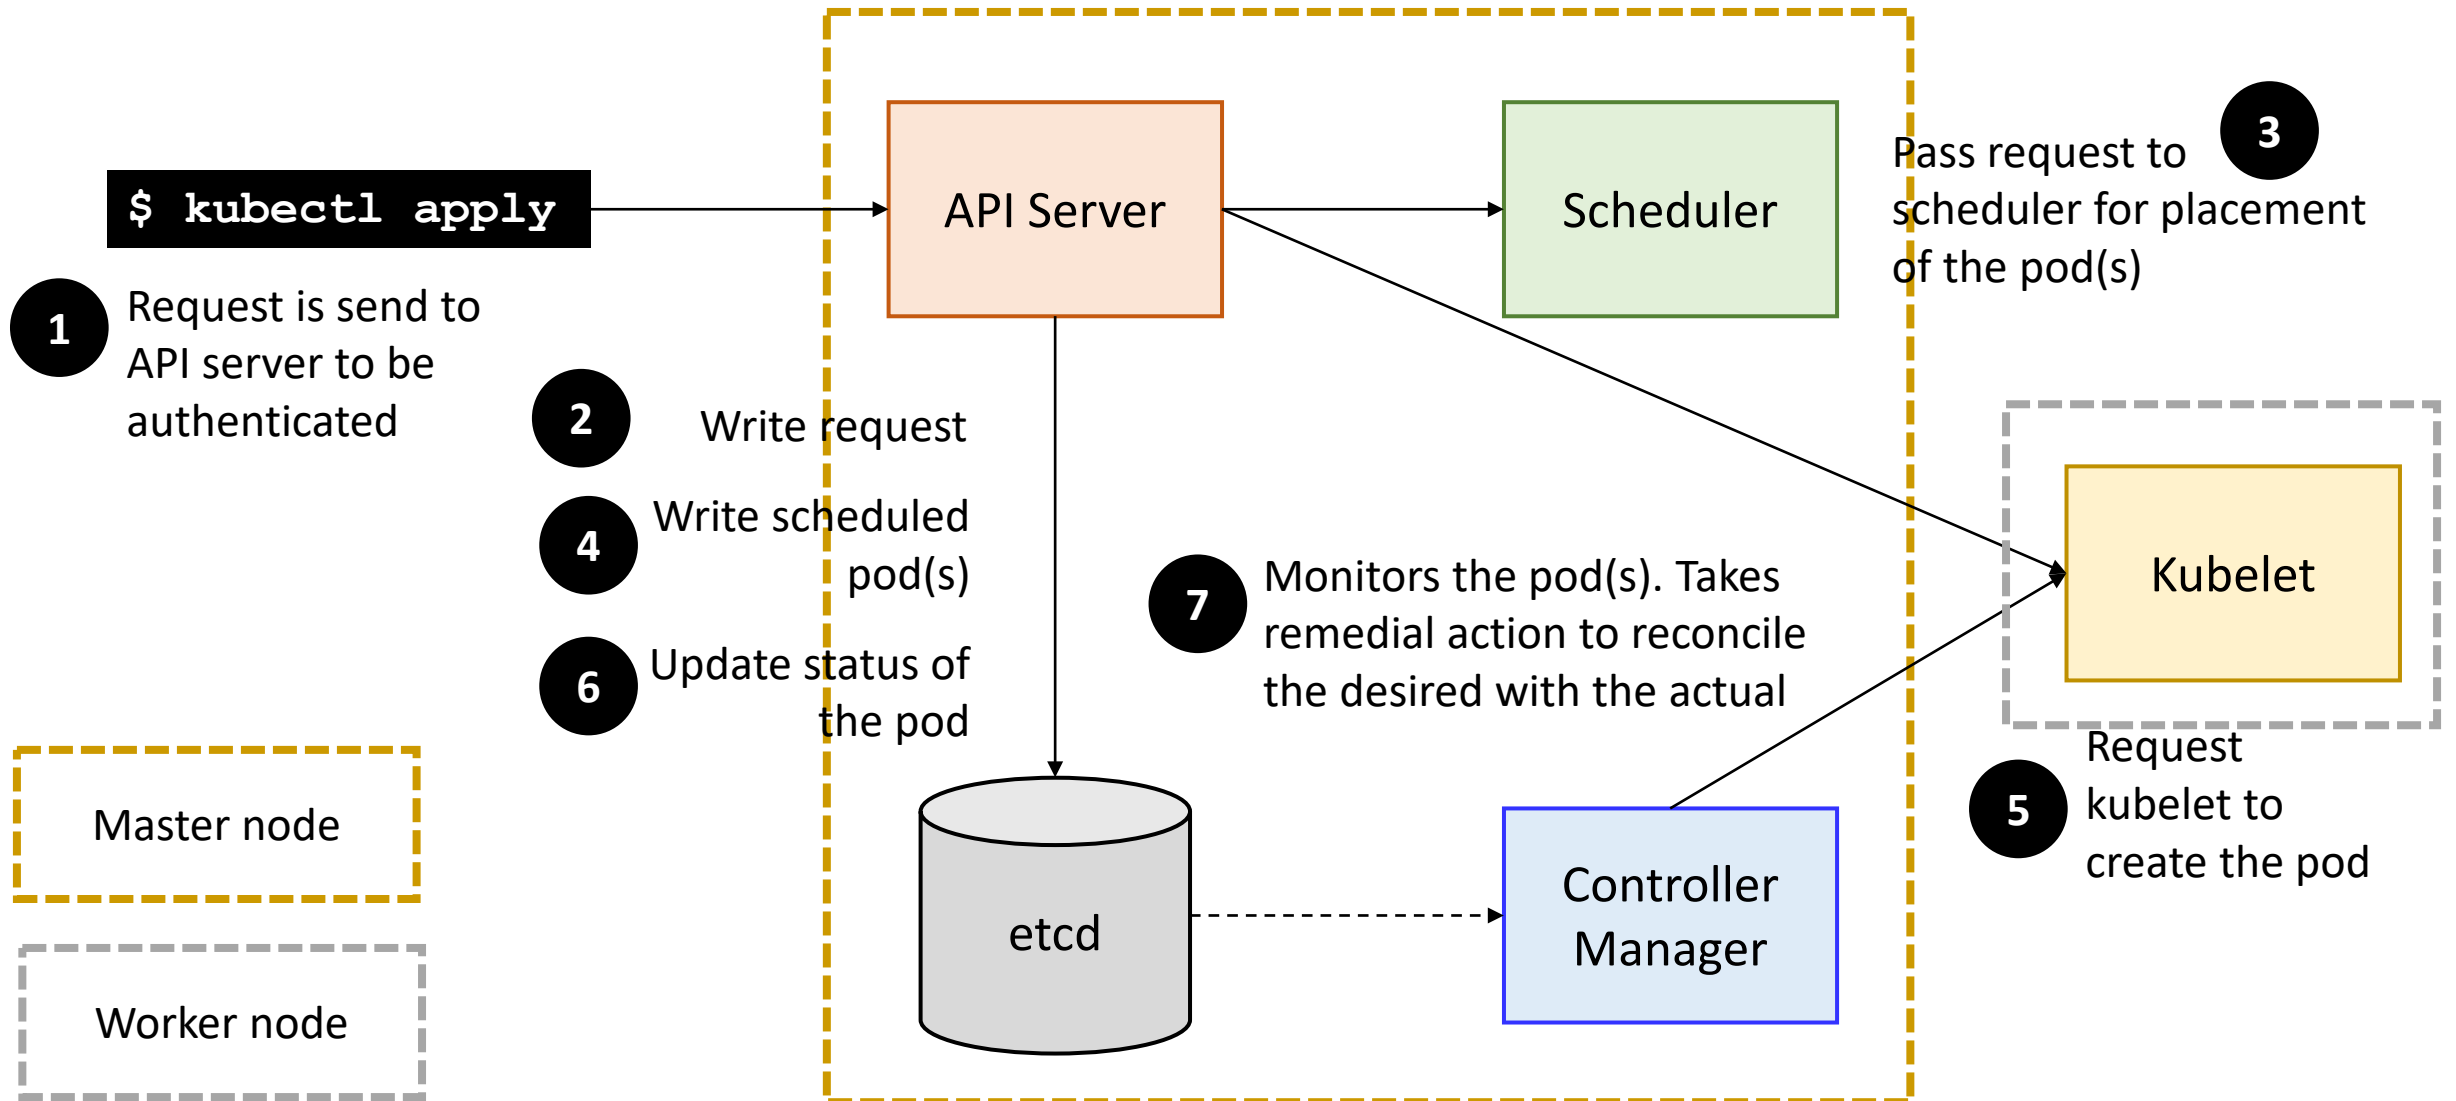
- Pod's container port is not exposed
- To access it need to bind the port to the host port

```
kubectl port-forward pod/myapp-pod 8080:3000
```

- Forwards traffic from port 8080 to Pods' port 3000
- Not a good way to access the application
 - Use for testing




How Kubernetes Provision Resources





Health Checks - Readiness


- Discover when a pod is ready to serve traffic
 - Eg. at startup - Pod is creating database
 - Eg. at steady state - scaling a large image and will not receive other request until the current one completes
- Will not route traffic to it until the pod is ready again. Traffic will be rerouted to other Pods





Health Checks - Liveness

- Checks if a Pod is dead or alive
- If Pod is dead, then remove the Pod and starts a new Pod to replace it
- Difference between liveness and readiness is that a Pod can fail readiness but is alive





Defining Probes

```
apiVersion: v1
```

```
...
```

```
spec:
```

```
  containers:
```

- name: myapp
 image: myapp:sha256:fadfec....
 imagePullPolicy: Always
 ports:

- name: app-port
 containerPort: 3000

```
  readinessProbe:
```

```
    httpGet:
```

```
      path: /ready  
      port: app-port  
      timeoutSeconds: 5  
      failureThreshold: 1
```

If these request returns a
status code of greater than 400
then it is consider a failure


```
  livenessProbe:
```

```
    httpGet:
```

```
      path: /  
      port: app-port  
      timeoutSeconds: 5  
      failureThreshold: 3  
      successThreshold: 1
```



Kubernetes





Deployments

- Almost always need more than a single Pod in production
- Deployments are used to create and deploy one or more Pods
- Deployment consist of
 - The number of Pods in the initial deployment
 - A template of the Pod which includes the Docker image, container port, etc.



Defining a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    matchLabels:
```

```
      name: myapp
```

```
  template:
```

```
    template:
```

```
      metadata:
```

```
        name: myapp-pod
```

```
        labels:
```

```
          name: myapp
```

```
      spec:
```

```
        containers:
```

```
          - name: myapp
```

```
            image: myapp@sha256:...
```

```
            imagePullPolicy: Always
```

```
            ports:
```

```
              - containerPort: 3000
```

Number of instances
in deployment

Pod definition

Criteria to identify the pods
belonging to this deployment



Deployment Management

- **Create a deployment**

```
kubectl apply -f deployment.yml
```

- **View all deployments**

```
kubectl get deploy -o wide
```

```
kubectl get deploy -o yaml
```

- **Detail information about a deployments**

```
kubectl describe deploy myapp-pod
```

- **Delete a deployment**

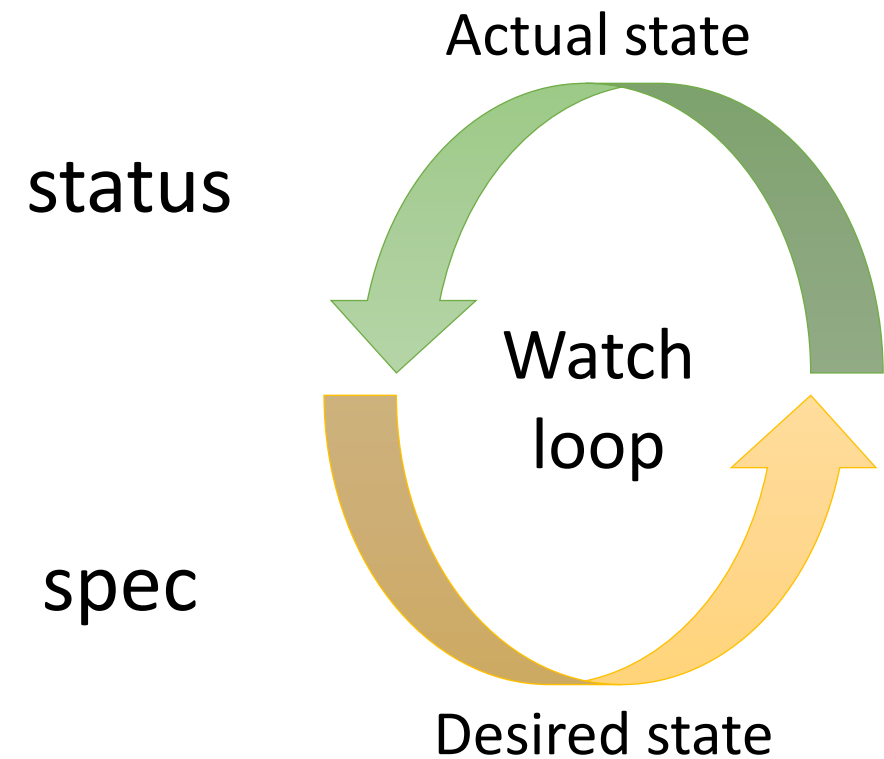
```
kubectl delete -f deployment.yml
```

```
kubectl delete deploy myapp-deployment
```



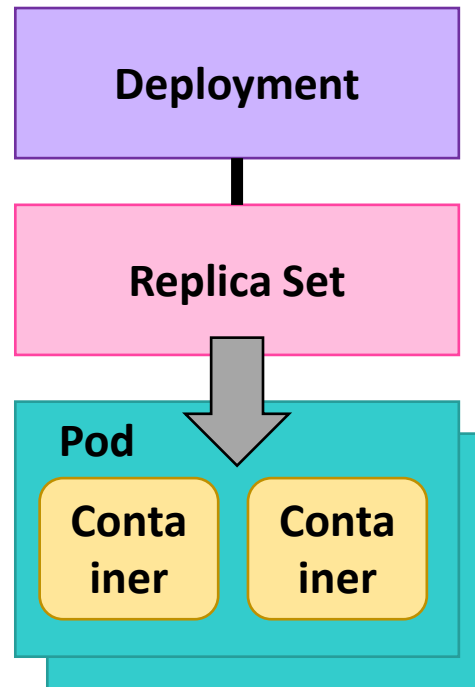
Replica Sets

- Deployments uses replica sets to manage Pods
- Replica sets ensures that the desired number of Pods are running
 - As defined in the `replica` attribute
- Kubernetes will match the actual state against the desired state
 - If the actual number of instances is less than the specification, Kubernetes will provision additional Pods so that the actual matches the desired





Kubernetes





Passing Values into Containers

- Configurations can be passed into containers in a Pod as configuration maps and/or secrets
 - Key/value pair files
- Difference between ConfigMap and Secret is that the latter values are base 64 encoded
- Passed into the container as
 - Environment variables viz. bind the values to environment variables
 - Mounted as a volume into a container




ConfigMap and Secrets

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: myapp-config
data:
  db_name: northwind
  db_host: myserver
  db_port: 3306
```

```
apiVersion: v1
kind: Secret
metadata:
  name: myapp-secret
data:
  db_user: ZnJlZA==
  db_password: eWFhYWRhYmFkb28=
```

```
echo -n 'fred' | base64
```



Encode value to base64



Injecting ConfigMap and Secret into Containers

```
containers:  
- name: myapp  
  image: myapp@sha256:...  
  ...  
  env:
```

```
- name: DB_HOST
```

```
  valueFrom:
```

```
    configMapKeyRef:
```

```
      name: myapp-config
```

```
      key: db_host
```

```
- name: DB_PASSWORD
```

```
  valueFrom:
```

```
    secretKeyRef:
```

```
      name: myapp-secret
```

```
      key: db_password
```

Bind the environment variable
DB_HOST to the following value

ConfigMap name

Key from ConfigMap


Secret name

Key from Secret

Bind the environment variable
DB_PASSWORD to the following
value




Kubernetes



Accessing the Application



A deployment





Pods are Ephemeral

- Pods are ephemeral
 - Can be reschedule to another node by the scheduler
 - Eg. when there is a node or network failure
- Clients cannot connect to pods directly via node
- Service provides a stable IP to the client
 - Acts as a proxy to a set of pods
 - Service keeps track of pods so clients connecting to them don't have to
- When pods are reschedule to another node the service is responsible for redirecting the request to another Pod instances





Service

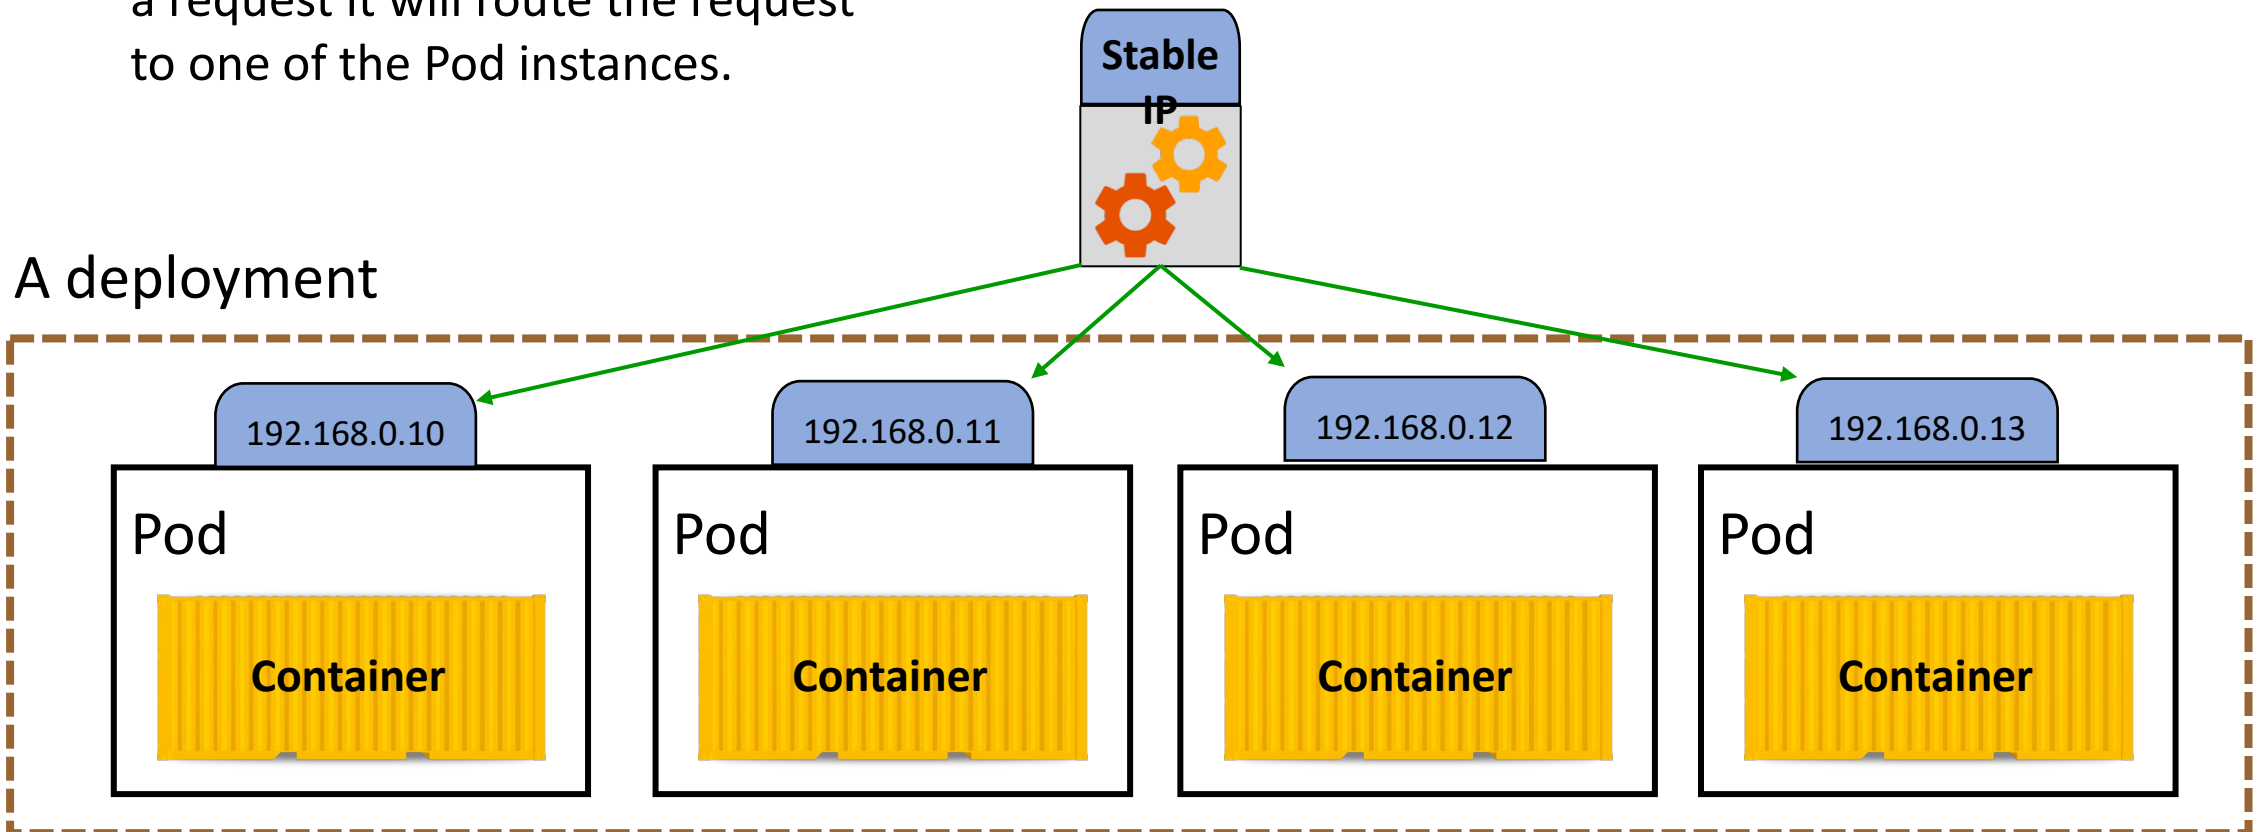
- Well known endpoint for a set of Pods in a deployment
- Services will route a request to a Pod under that the service controls
 - Also provides simple load balancing
- Pods are selected to be in a service based on their labels
 - For a Pod to be in a service, the Pod only has to match some of its Pod labels
 - Pods and services are loosely coupled
- Services are durable, unlike Pods which are ephemeral
 - Static IP address
 - Static namespace DNS name

Accessing the Application

Provides a durable IP address to the client. When a service receives a request it will route the request to one of the Pod instances.



Endpoint keeps track of the Pod's IP address in a deployment





Defining a Service

```
apiVersion: v1
kind: Service
metadata:
  name: myapp
```

Service name

```
spec:
  type: ClusterIP
  selector:
```

Specify the type of service that is exposed

```
    name: myapp
    version: v1
```

Route service to Pods that matches these labels

```
  ports:
```

```
    - name: http:
      port: 8080
      targetPort: 3000
      protocol: TCP
```

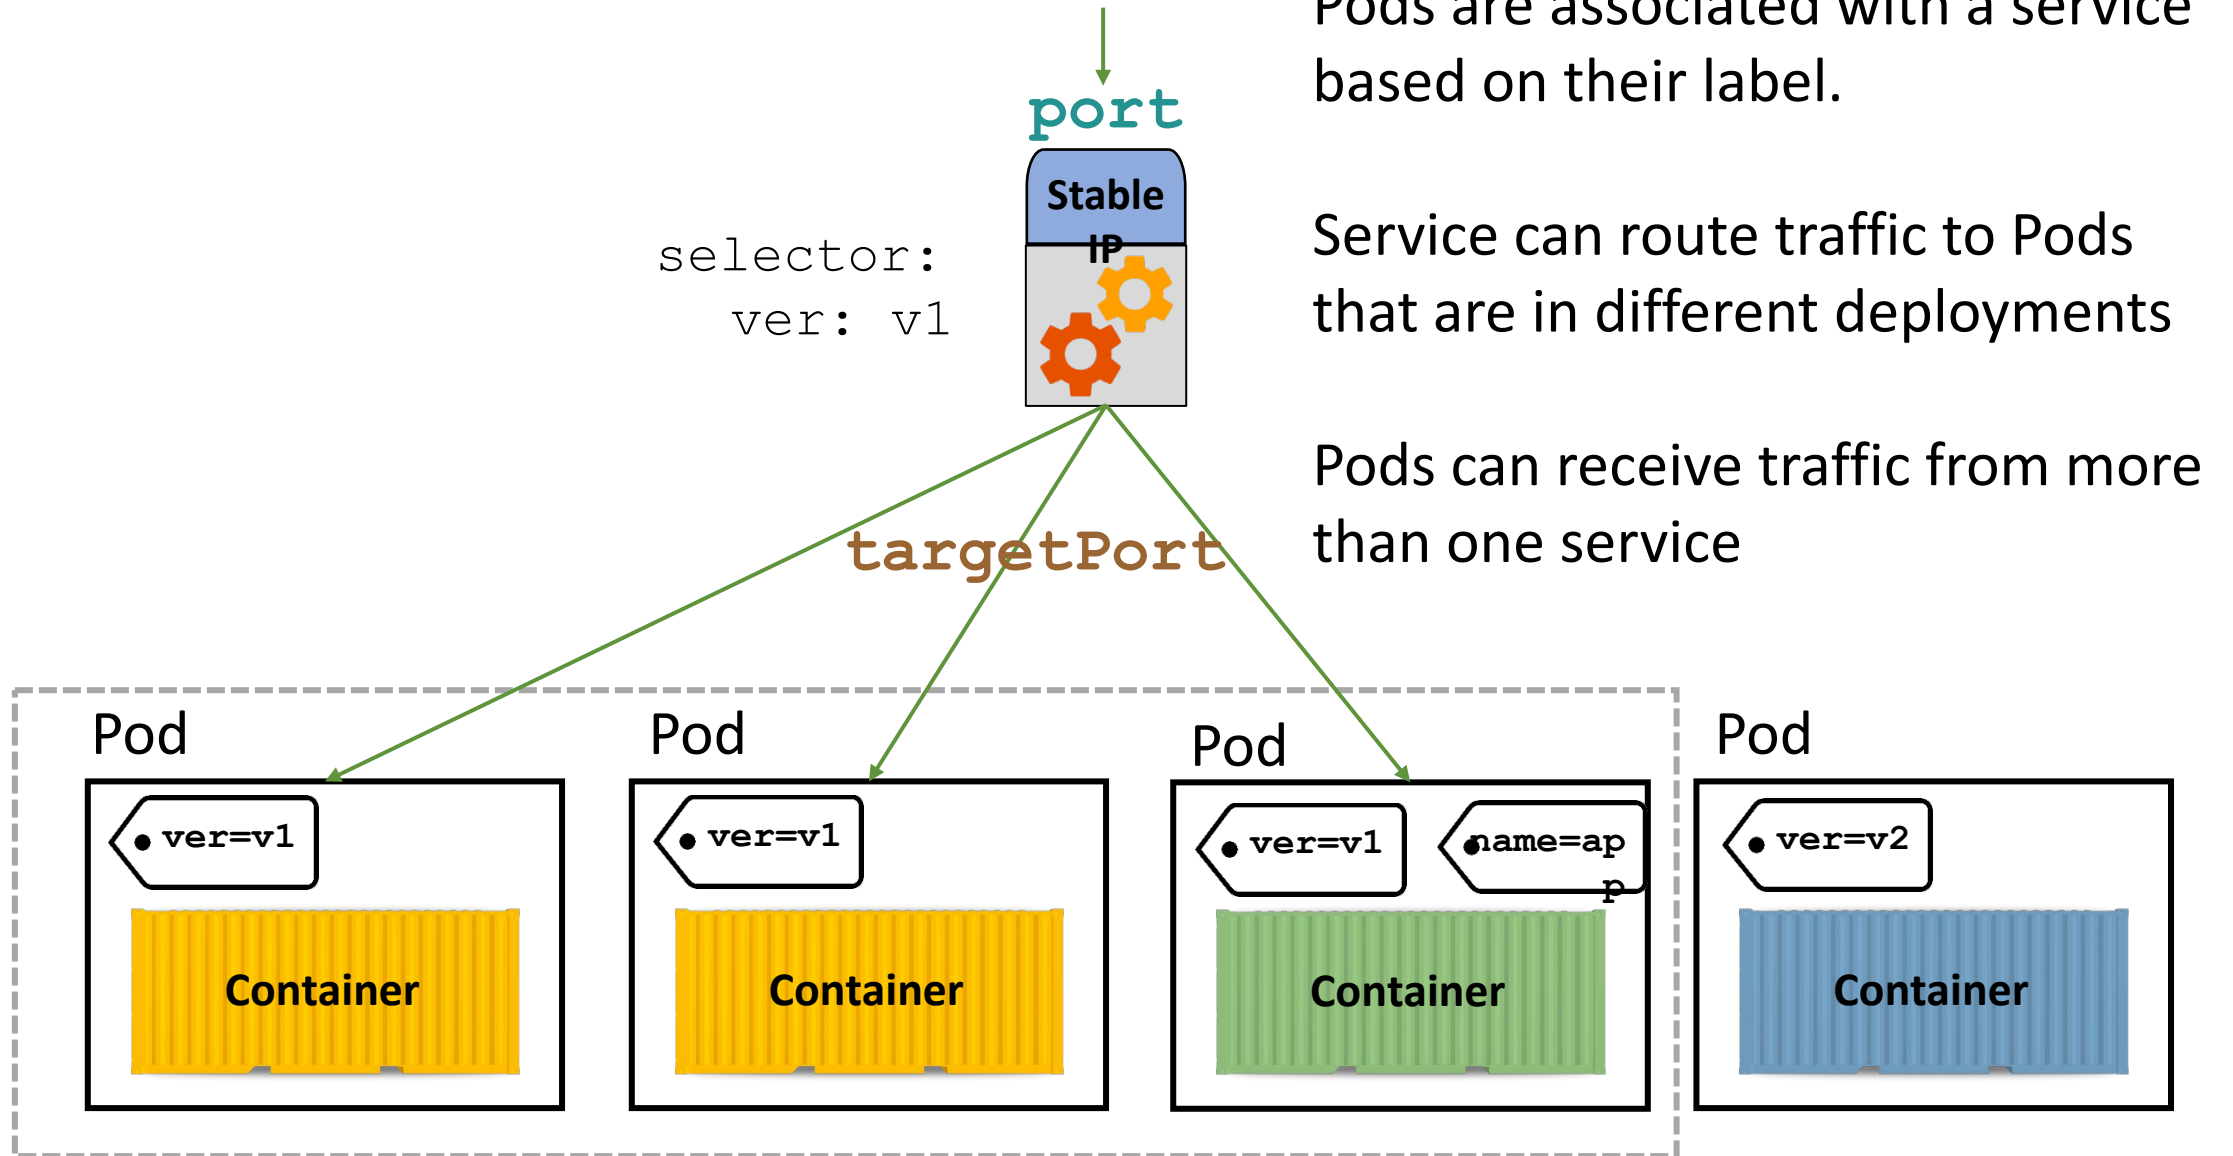
The port(s) that are exposed by this service

The port(s) that are exposed by this service

Route traffic from port (8080) to the Pod's port (targetPort 3000)



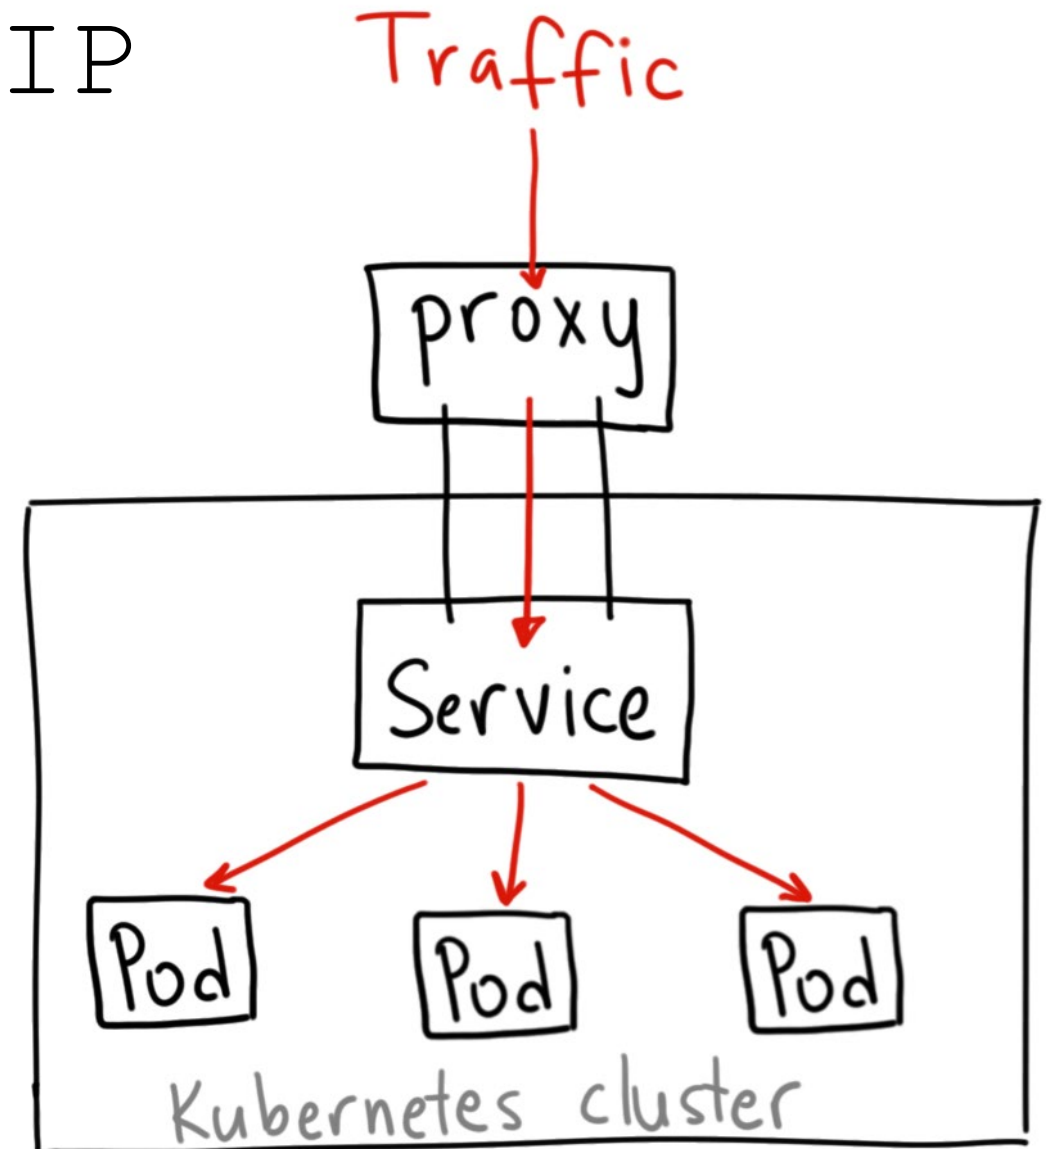
Services and Pods





Service Type - ClusterIP

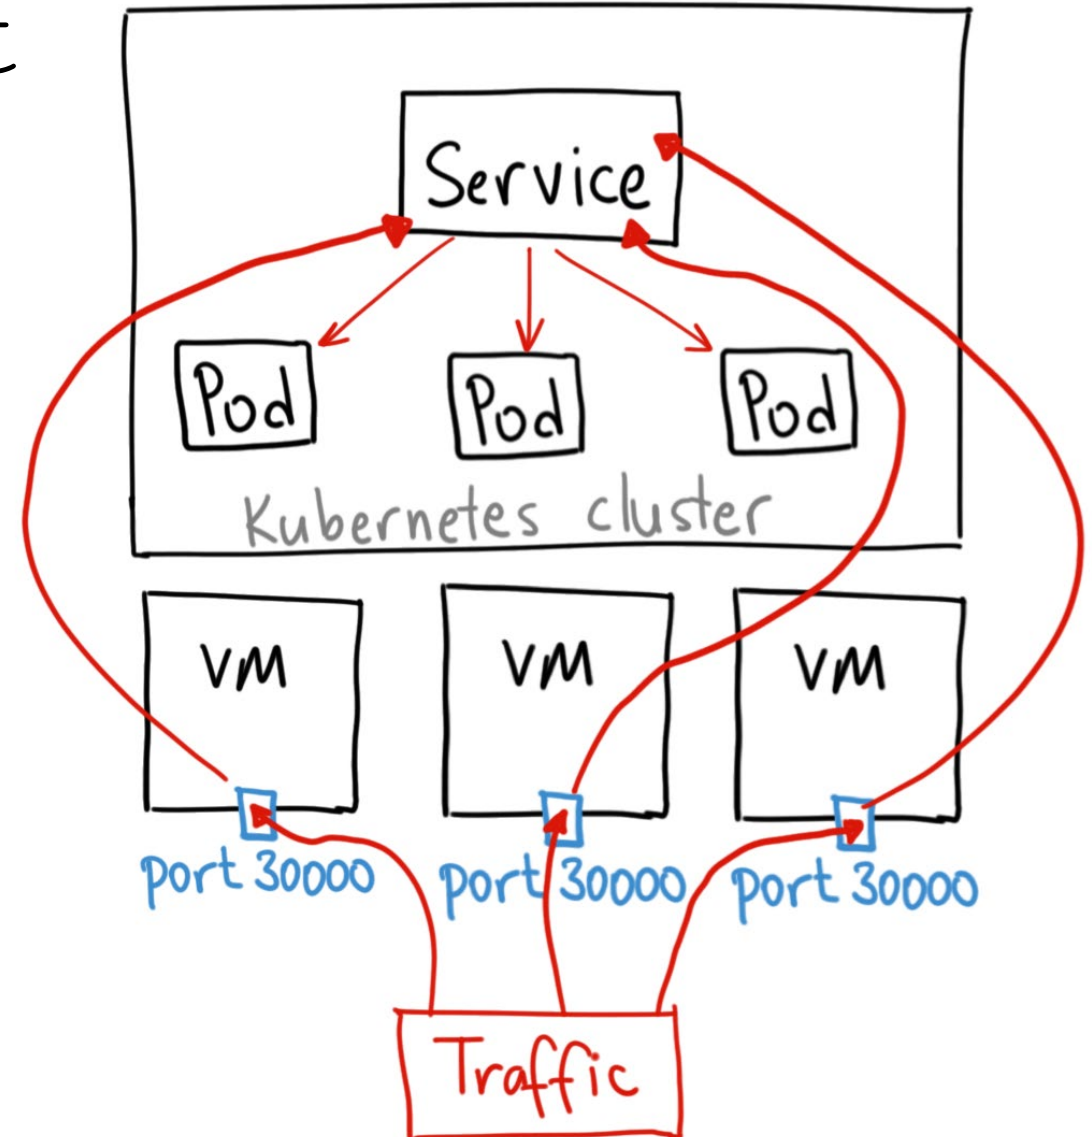
- Provision service IP address inside the cluster
- The IP address is not accessible from outside of the cluster
- This is the default





Service Type - NodePort


- Opens a port, 30000 in this case, in all the node
- Traffic arriving at port 30000 on any of the nodes will be routed to the service
 - Expose as `ClusterIP`
- Make services appear local viz. the service can be accessed with localhost
 - Extra hop if the node does not have to Pod
- Need to pick a cluster node and the exposed port (node port) to access the service






Service Type - LoadBalancer

- A load balancer will route traffic to the service
- Traffic coming from the load balancer will be distributed to a node according to its routing policy
 - Exposed as `NodePort`
- The load balancer is accessible from outside of the cluster
- `LoadBalancer` service type will be provisioned by the underlying cloud platform
 - By the cloud controller manager





Service





Pod to Service Communication

- Kubernetes creates an entry in its internal DNS (KubeDNS)
`<service_name>.<namespace>.svc.cluster.local`
- A pod can access a service either with
 - service name if pods and service are in the same namespace
 - FQDN if service and pods are in different namespaces



Accessing the Service

- May need to access the service from outside the cluster
 - For testing

- Forward traffic from the host into the cluster's IP

```
kubectl port-forward svc/<service_name> 8080:3000
```

- Port map 8080 from the host to port 3000 exposed by the service

- Or start a kube-proxy

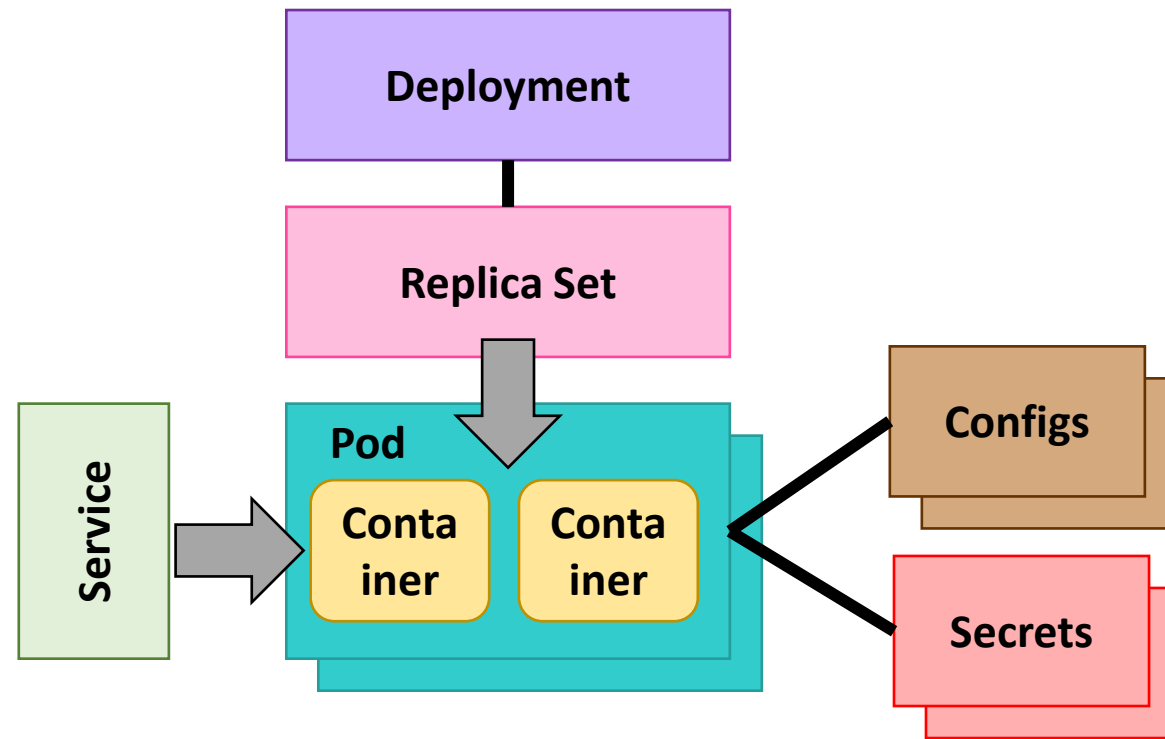
```
kubectl proxy --port=8080
```

- Access the service with the following URL

```
http://localhost:8080/api/v1/namespaces/<namespace>/services/  
http:<service_name>:3000/proxy/
```



Kubernetes






Rolling Updates

- Applications need to be updated
 - Pods with new images
- Rolling updates allow deployments to be updated without any downtime
 - Replace old Pods with the new gradually
 - When old Pods are no longer serving request
- Alternative to rolling updates is 'recreate'
 - Kill all existing Pods before creating new ones




Rolling Deployment vs Recreate

Rolling Upgrade




Recreate





Rolling Update Illustrated

maxSurge: 1
maxUnavailable: 0





Defining a Rolling Update

Number of seconds for
Kubernetes to wait for
the application to be
ready

```
apiVersion: apps/v1
kind: Deployment
...
spec:
```

```
  replicas: 3
  minReadySeconds: 5
```

Use the rolling update
strategy for this
deployment with 3
replicas

```
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
```

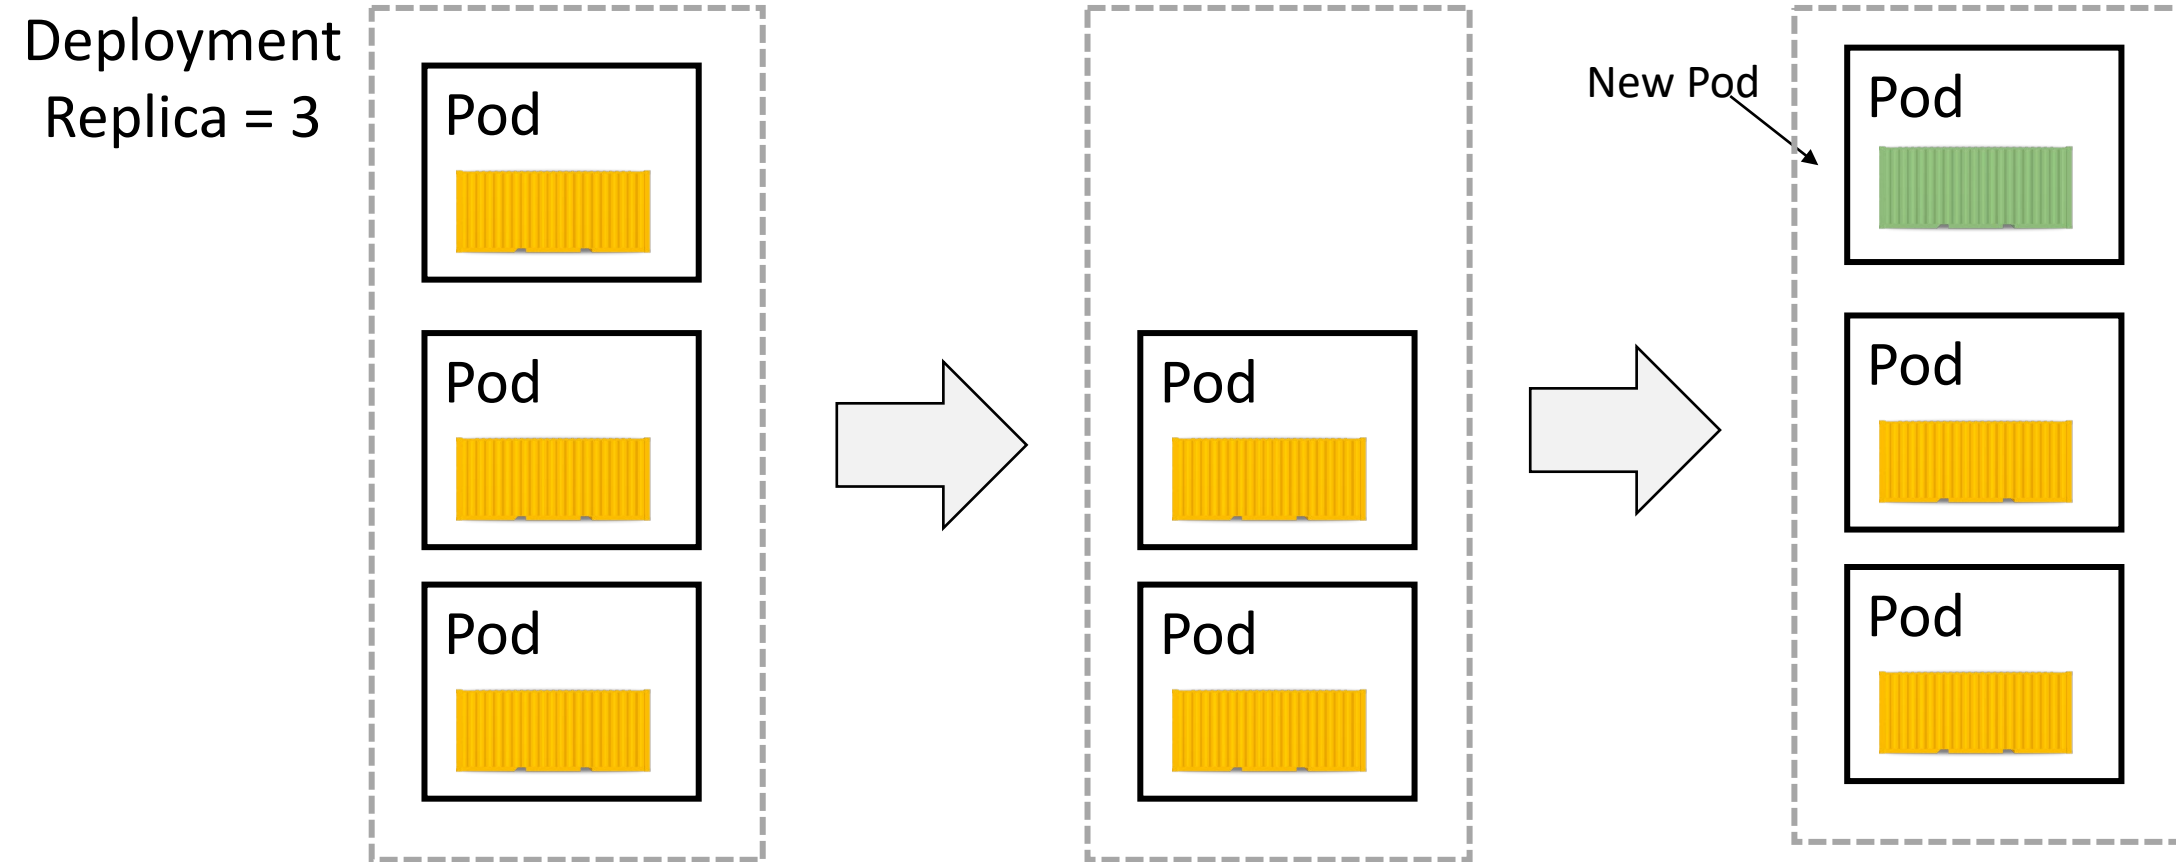
The policy for updating the deployment
is controlled by `maxSurge` and
`maxUnavailable`

Pods in this deployment will be updated
one at a time (`maxSurge`). At any time
the number of Pods will not fall below 3
(`maxUnavailable`)



Rolling Update Illustrated


maxSurge: 0
maxUnavailable: 1




```
kubectl apply -f updated_deployment.yml
```



Rolling Updates



```
replicas: 5  
maxSurge: 5  
maxUnavailable: 0
```



```
replicas: 5  
maxSurge: 1  
maxUnavailable: 1
```



Rollback

```
kubectl apply -f dep-v1.yml
```

```
kubectl apply -f svc.yml
```

```
kubectl apply -f dep-v2.yml
```

```
kubectl rollout history deployment myapp-deployment
```

REVISION	CHANGE-CAUSE
----------	--------------

1	
---	--

2	
---	--

```
kubectl rollout undo deployment myapp-deployment --to-revision=1
```



Managing Updates

- **Apply an update**

```
kubectl apply -f deployment-next.yml
```

- **Check update status**

```
kubectl rollout status deployment <deployment_name>
```

- **See the revision history of the deployment**

```
kubectl rollout history deployment <deployment_name>
```



- **Rollback to a previous version**

```
kubectl rollout undo deployment <deployment_name> \  
--to-revision=<rev>
```



Blue Green Deployment

- Release model where 2 versions of an application is deployed side-by-side
 - Old version is called blue, the new version is called green
- When the green deployment is ready to receive traffic, reconfigure load balancer to forward request from blue to green
- Blue is maintained for a period before decommissioning
 - For rollback if there are issues with the green deployment





Example - Blue Green Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deploy-v1
  namespace: prod
spec:
  replicas: 3
  selector:
    matchLabels:
      name: myapp-po
      version: v1
  ...
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deploy-v2
  namespace: prod
spec:
  replicas: 3
  selector:
    matchLabels:
      name: myapp-po
      version: v2
  ...
```

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-svc
  namespace: prod
spec:
  type: ClusterIP
  selector:
    version: v1
```

Change Service selector from v1 to v2 when v2 pods are ready

```
kubectl get po -n prod
-1 name=myapp-po,version=v2 \
-o custom-columns='NAME:.metadata.name,READY:.status.conditions[?(@type=="Ready")].status'
```

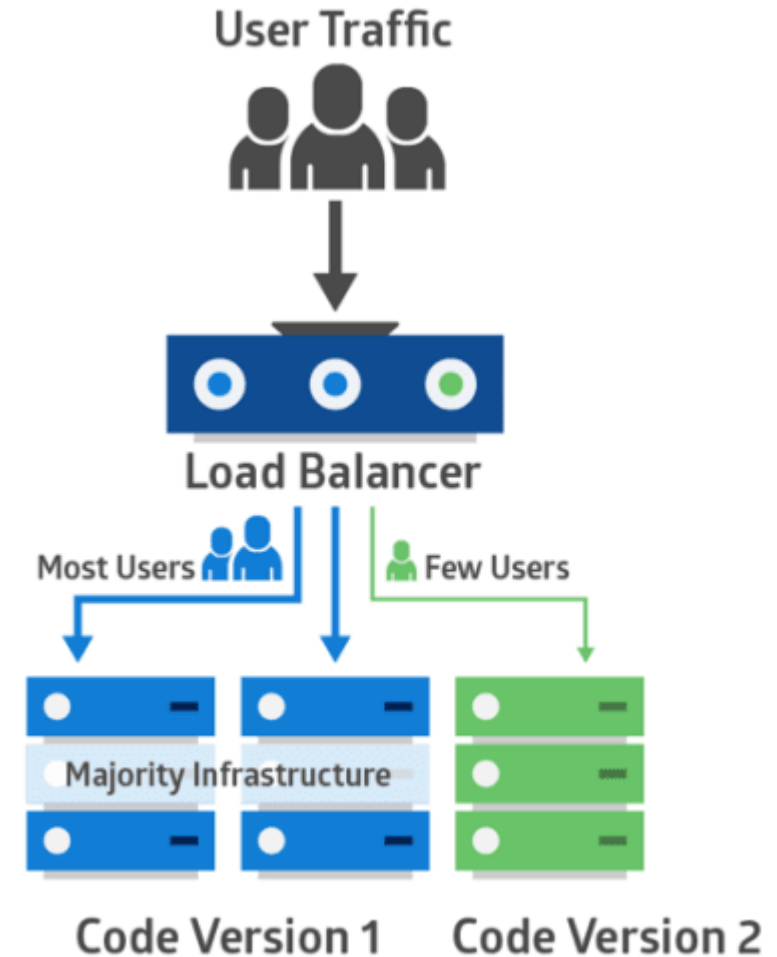
Select pods with these labels only

Custom column for displaying pod conditions



Canary Deployment

- Making staged deployment where new releases is tested with a small subset of users
 - Part of request, e.g. 30% is served by the canary
- Objective of canary is to
 - Solicit feedback from users
 - Test the new release in production, if there are issues, will only affect a small percentage of the users





Canary Deployment

- Vanilla Kubernetes does not support canary deployment
- Use multiple deployment to mix different versions of the same application
 - Assign a common label to both deployments
 - Service forwards traffic to pods by selecting the common label
 - If the split is 1 in 4, 25% of the request goes to the new application, then deploy 1 new pod and 3 old pods
- Other methods of creating canary deployment
 - Ingress Nginx proprietary feature for marking an endpoint as canary
 - Istio using virtual service to split the traffic between different services



Example – Canary Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deploy-v1
  namespace: prod
spec:
  replicas: 3
  selector:
    matchLabels:
      name: myapp-po
      version: v1
  ...
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp-deploy-v2
  namespace: prod
spec:
  replicas: 1
  selector:
    matchLabels:
      name: myapp-po
      version: v2
  ...
```

Common label on both
the deployments

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-svc
  namespace: prod
spec:
  type:
  selector: ClusterIP
  version: myapp-po
```

Use the service to randomly route
incoming traffic to both the
deployments using the common
label(s)



Appendix



Architecture

