



Tools for Working with Kubernetes




Kustomize



Multiple Environments

- Different environments will require different tweaks to the manifests
 - Different configurations
 - Some can be configured with `ConfigMaps` or `Secrets`
 - Some configurations cannot eg. replicas, image name
 - Different attributes
 - Require resource limits in production
 - Additional resources like cache in production





Using Kubernetes In Multiple Environments



- DRY – Don't Repeat Yourself
- Develop a Helm chart
- Nontrivial
 - Need to develop templates that cover all possible scenarios
 - Learn Go templates, template functions, conventions, etc
 - Debugging!
- Significant time investment
- Helm provides
 - Structured way of customizing deployments
 - Deployment management - install, delete , upgrades



- WET – Write Everything Twice
- Duplicate and modify the original Kubernetes manifest
- Need to actively update downstream manifest when upstream changes
 - Eg. bug fixes, improvement, etc.
- Copy and paste nature is
 - Easy to use
 - No new concepts to learn
 - Fast



Kustomize

- Kustomize is a tool for customizing Kubernetes deployments
 - Apply a set of changes to a base Kubernetes manifest
 - Produces a new set of manifest
 - Original manifest (base) is untouched/unchanged
 - Template free viz. logic less
- Kustomize allow you to consume and customize upstream Kubernetes manifest without forking them
- Open source
 - Available as standalone - <https://kubectldocs.kubernetes.io/installation/kustomize/>
 - Integrated with kubectl (post 1.14) with 'apply -k'



Customizing Kubernetes Deployments

- Customize transform a set of related Kubernetes manifest by adding or replacing fields and objects
 - Can delete fields but not common
- Modifications are updates can be applied to
 - All resources - eg adding a common label
 - Specific resource by providing patterns that match the resource
- Input Kubernetes manifest are specified with the `resources` field



Example - Common Fields

kustomization.yaml

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- pod.yaml
commonLabels:
  env: stage
namePrefix: eng-
namespace: stage-ns
```

pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: app-pod
  labels:
    name: app
spec:
  containers:
  - name: app-container
    image: nwapp:v1
```

```
apiVersion: v1
kind: Pod
metadata:
  name: eng-app-pod
  labels:
    name: app
    env: eng
  namespace: stage-ns
```

```
spec:
  containers:
  - name: app-container
    image: nwapp:v1
```

With kustomize

```
kustomize build | kubectl apply -f -
```

With kubectl

```
kubectl apply -k .
```

Directory with
kustomization.yaml



Example - Deployments

```
kustomization.yaml
apiVersion: ...
kind: Kustomization
```

replicas:

```
- name: web-deploy
  count: 1
```

images:

```
- name: resilio/sync
  newName: eeacms/rsync
  newTag: 2.3
- name: nginx
  newTag: 1.21.1-perl
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
  name: web-deploy 1
```

```
spec:
```

```
  replicas: 2
```

```
  selector:
```

```
    ...
```

```
  template:
```

```
    ...
```

```
    spec:
```

```
      containers:
```

```
        - name: file-sync
```

```
          image: resilio/sync:2.7.2
```

```
        ...
```

```
        - name: nginx
```

```
          image: nginx:1.20.1
```

```
        ...
```

nginx:1.21.1-perl

eeacms/rsync:2.3



Patching

- Overrides or add fields to the target resource
- A list of patch files are provided
 - Order of patches applied is the order in which they are listed
- Patch file is similar to the resource to be patched
 - Should have `apiVersion`, `kind` and `name`; used by Kustomize to select the resource to be patched
 - The remainder of the patch file used to overlay on the resource viz. metadata and spec sections
 - Attributes with the same name are replaced
 - Additional attributes are added



Example - Patching

app.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: myapp

spec:

replicas: 2

selector:

...

template:

...

spec:

containers:

- name: nwapp

image: nwapp:v2

ports:

- containerPort: 3000

app.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: myapp

spec:

replicas: 2

selector:

...

template:

...

spec:

containers:

- name: nwapp

image: nwapp:v2

ports:

- containerPort: 3000

resources:


requests:

cpu: 100m

memory: 128Mi

Customized
resource

Add these





Example - Patching

kustomization.yaml

```
apiVersion:
kustomize.config.k8s.io/v1beta1
kind: Kustomization
```

```
resources:
- app.yaml
```

patchesStrategicMerge:

```
- patch.yaml
```

patch.yaml

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: myapp
```

```
spec:
```

```
  template:
```

```
    spec:
```

```
      containers:
```

```
        - name: nwapp
```


```
          resources:
```

```
            requests:
```

```
              cpu: 100m
```

```
              memory: 128Mi
```

Overlay, provide enough
context to patch the original

A black curved arrow originates from the text "Overlay, provide enough context to patch the original" and points towards the "resources:" section of the patch.yaml file.



ConfigMap and Secrets

- Can customize an existing configuration or secret specific to an environment
 - Modify an existing configuration by merging with new keys or replacing existing keys
 - Generate new configurations and secrets
- **ConfigMap and Secrets** behaviour
 - Pods that references values in these 2 resources are not recreated if the content changes eg. value of a secret key is replaced
 - Referenced with either `envFrom` or `valueFrom`
 - Pods that mount these 2 resources will be recreated



Example - Override Existing Configurations

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
```

```
configMapGenerator:
```

- name: **app-cm**
 behavior: merge
 literals:
 - DB_USER=barney
 - TRACE_ENABLE="1"

Add or update the
values in the base

```
secretGenerator:
```

- name: **app-secret**
 behavior: merge
 type: Opaque
 literals:
 - DB_PASSWORD=mypassword

Replaced and
added to the base

Replace the base

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-cm
data:
  DB_HOST: localhost
  DB_PORT: 3306
  DB_USER: fred
  DB_NAME: inventory
```

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
type: Opaque
data:
  DB_PASSWORD: bX1zZWNYZ
  XQ=
```



Example - Generating ConfigMap and Secret

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Config
```

```
resources:
- app.yaml
```

generatorOptions:

```
  labels:
    env: stage
```

configMapGenerator:

```
- name: app-cm
  behavior: create
  literals:
  - DB_HOST=dbserver
    DB_USER=barney
```

Add labels to all the
ConfigMap and Secret

Create a ConfigMap
call app-cm

One or me
configurations

app.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
```

```
  ...
spec:
```

```
  ...
  containers:
  - name: app-container
    image: nwapp:v1
    envFrom:
      configMapRef
        name: app-cm
```

Reference a non
existence ConfigMap

Similar for secretGenerator



Example - Generating ConfigMap and Secret

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-cm-bfA9Ekh0c5
  labels:
    env: stage
data:
  DB_HOST: dbserver
  DB_USER: barney
```

Generated ConfigMap with a unique suffix.

A new suffix is generated everytime kustomize runs

Uses this ConfigMap when it is referenced from containers

Suffix hash forces Kubernetes to recreate the pods when configurations and secrets are updated

```
app.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  ...
spec:
  ...
  containers:
    - name: app-container
      image: nwapp:v1
      envFrom:
        - app-cm-bfA9Ekh0c5
```



Replacements

- Need to know certain info from resources generated by Kustomize
 - Eg. generated a service for a database with a prefix. Need to pass this service name to a deployment that uses the database
- Used to copy fields from one source to one or more targets
 - Eg copying the service name to environment variables
- Source is the resource that will provide the value; must be a single resource
- Target(s) are the resources that the fields will be replaced
 - Include and reject sets
- Use `fieldPath` to select specific fields as source and targets
 - See <https://kubectldocs.kubernetes.io/references/kustomize/kustomization/replacements/#delimiter>



Example - Replacement

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
```

replacements:

```
- source:
  kind: Service
  fieldPath: metadata.name
```

targets:

```
- select:
  kind: Deployment
  name: app
  fieldPaths:
  - spec.template.spec.containers.[name=myapp].env.[name=DB_HOST].value
```

Path to the value to be replaced

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app
spec:
  template:
    spec:
      containers:
      - name: myapp
        env:
        - name: DB_HOST
          value: old_value
```


Selects this resource

Update the value
of this field



Directory Layout

- Create a structure to support customizing a set of Kubernetes manifest
 - Support multiple environments (eg dev, test, prod) without interfering with another environment
 - Any changes will be picked up by all upstreams viz. dev, test, prod, etc
- Base manifest
 - Generic Kubernetes manifest applicable to all environments
- Overlays
 - Customize configurations to be applied to the base manifest
 - One per environment





Directory Structure

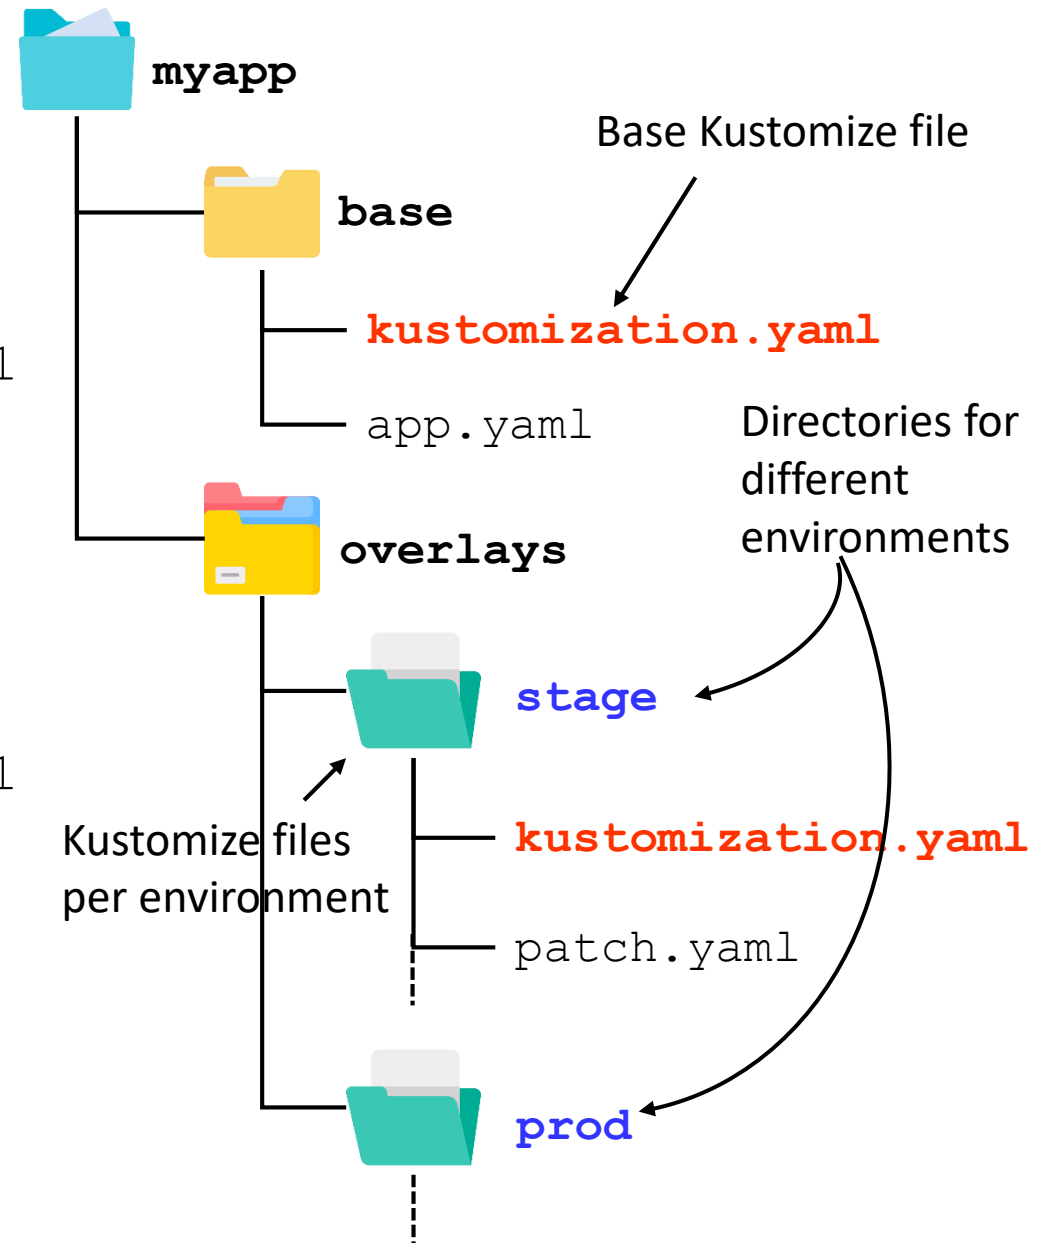
base/kustomization.yaml

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- app.yaml
```

overlays/stage/kustomization.yaml

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- ../../base
```

```
# site specific customization
namePrefix: ...
```

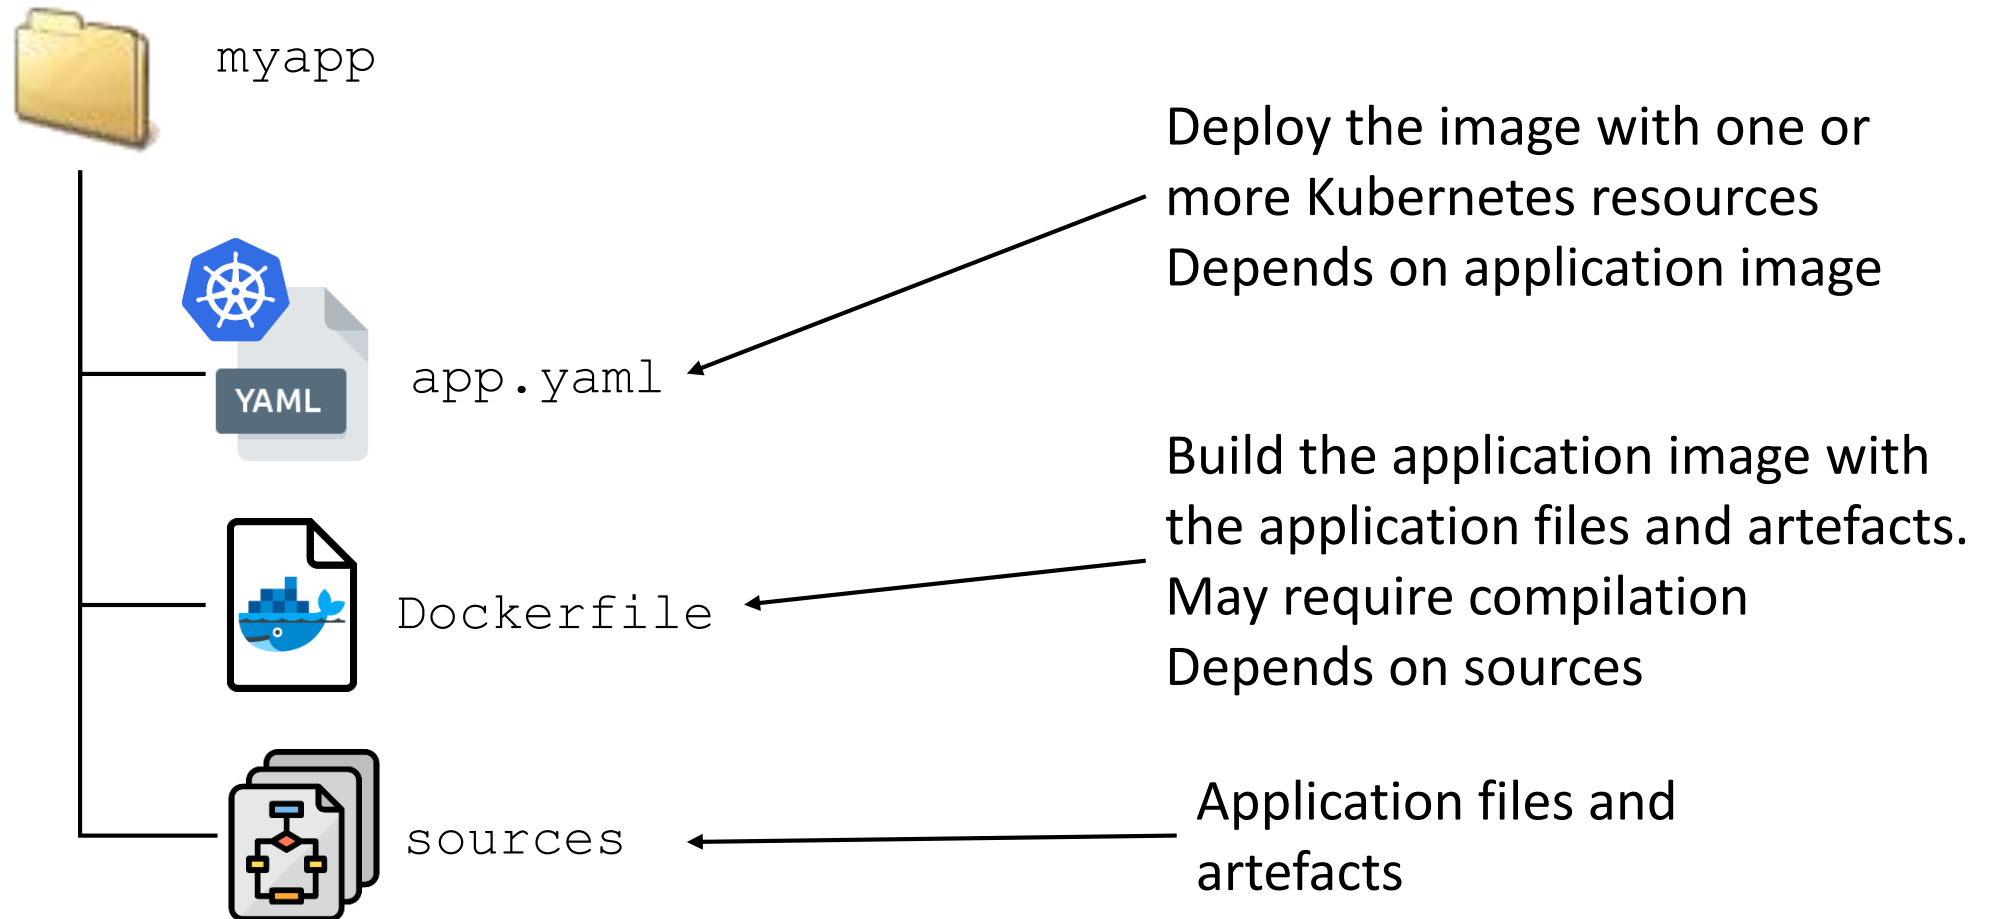




Skaffold

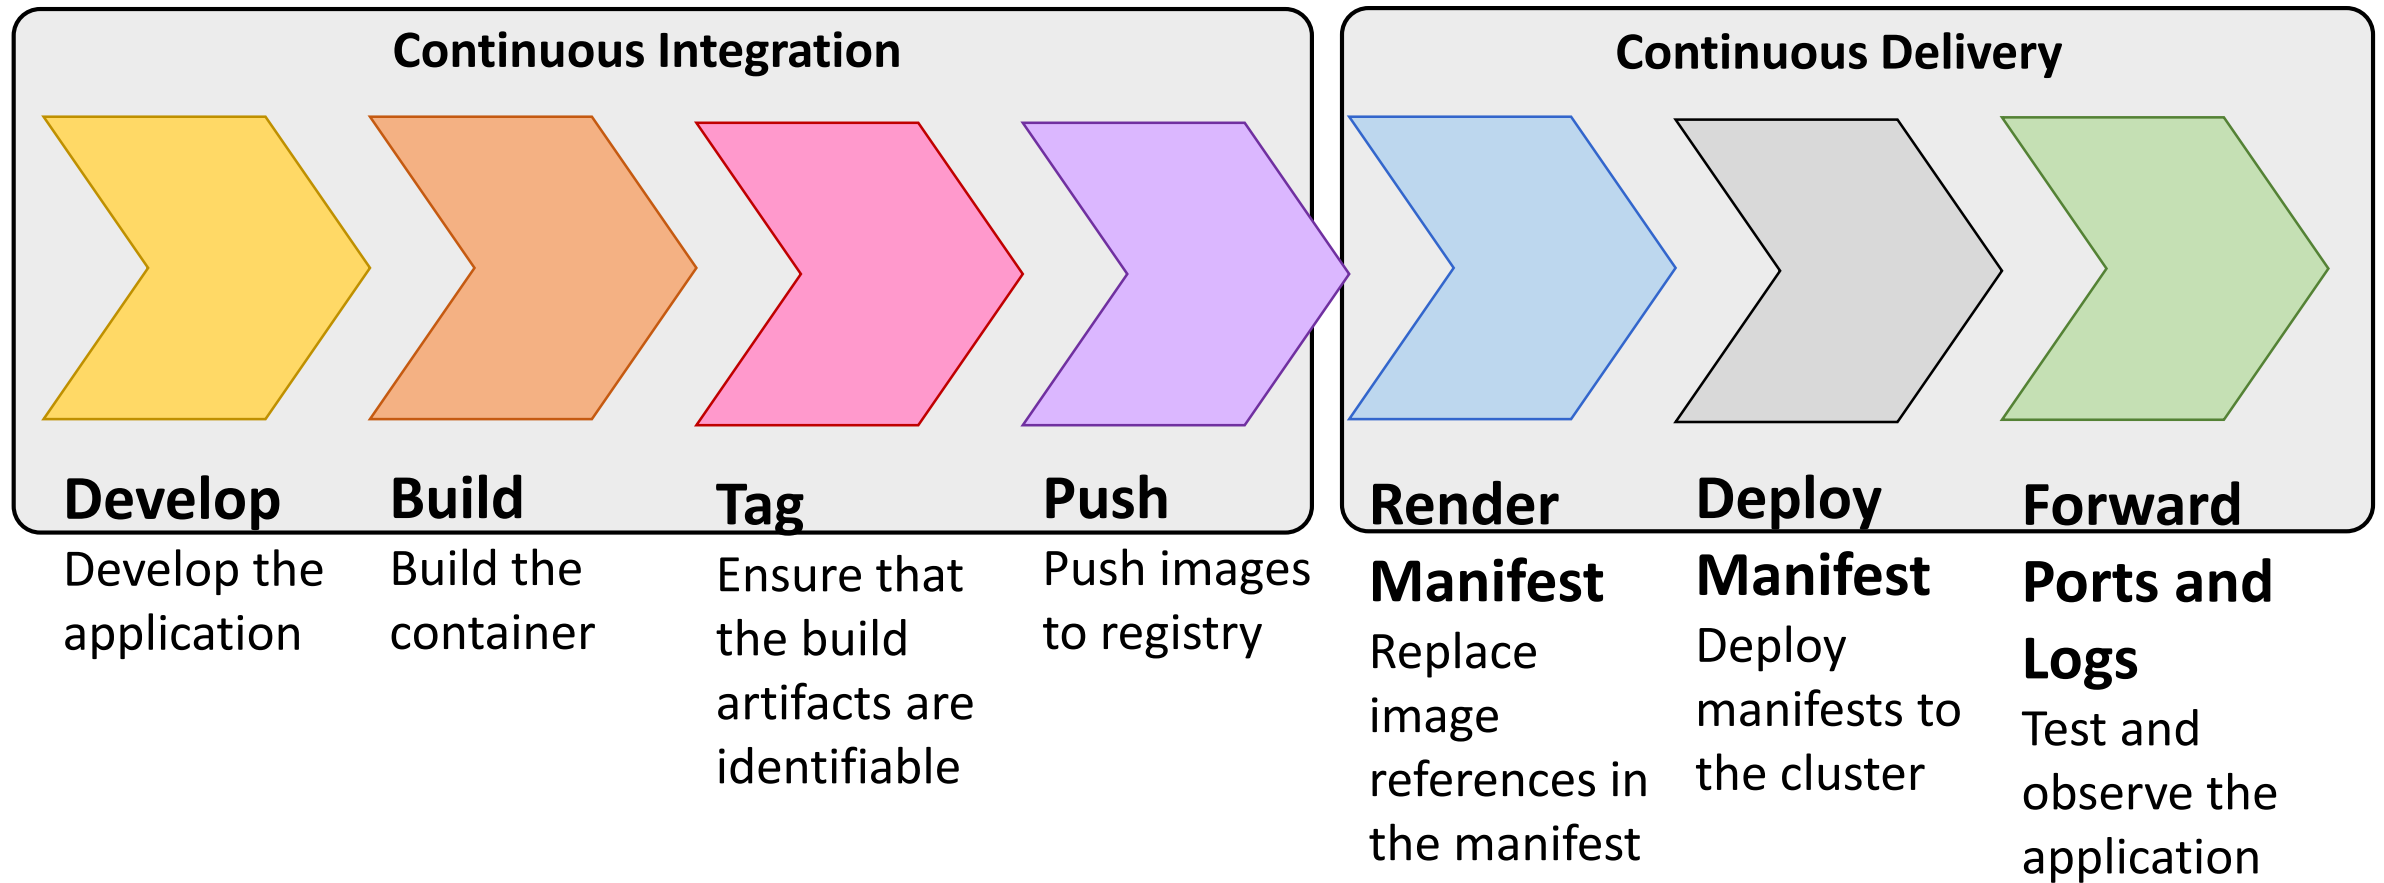


Kubernetes Application Setup





Developing, Debugging and Deploying



Repeat the process if there is any errors or bugs in any of the steps

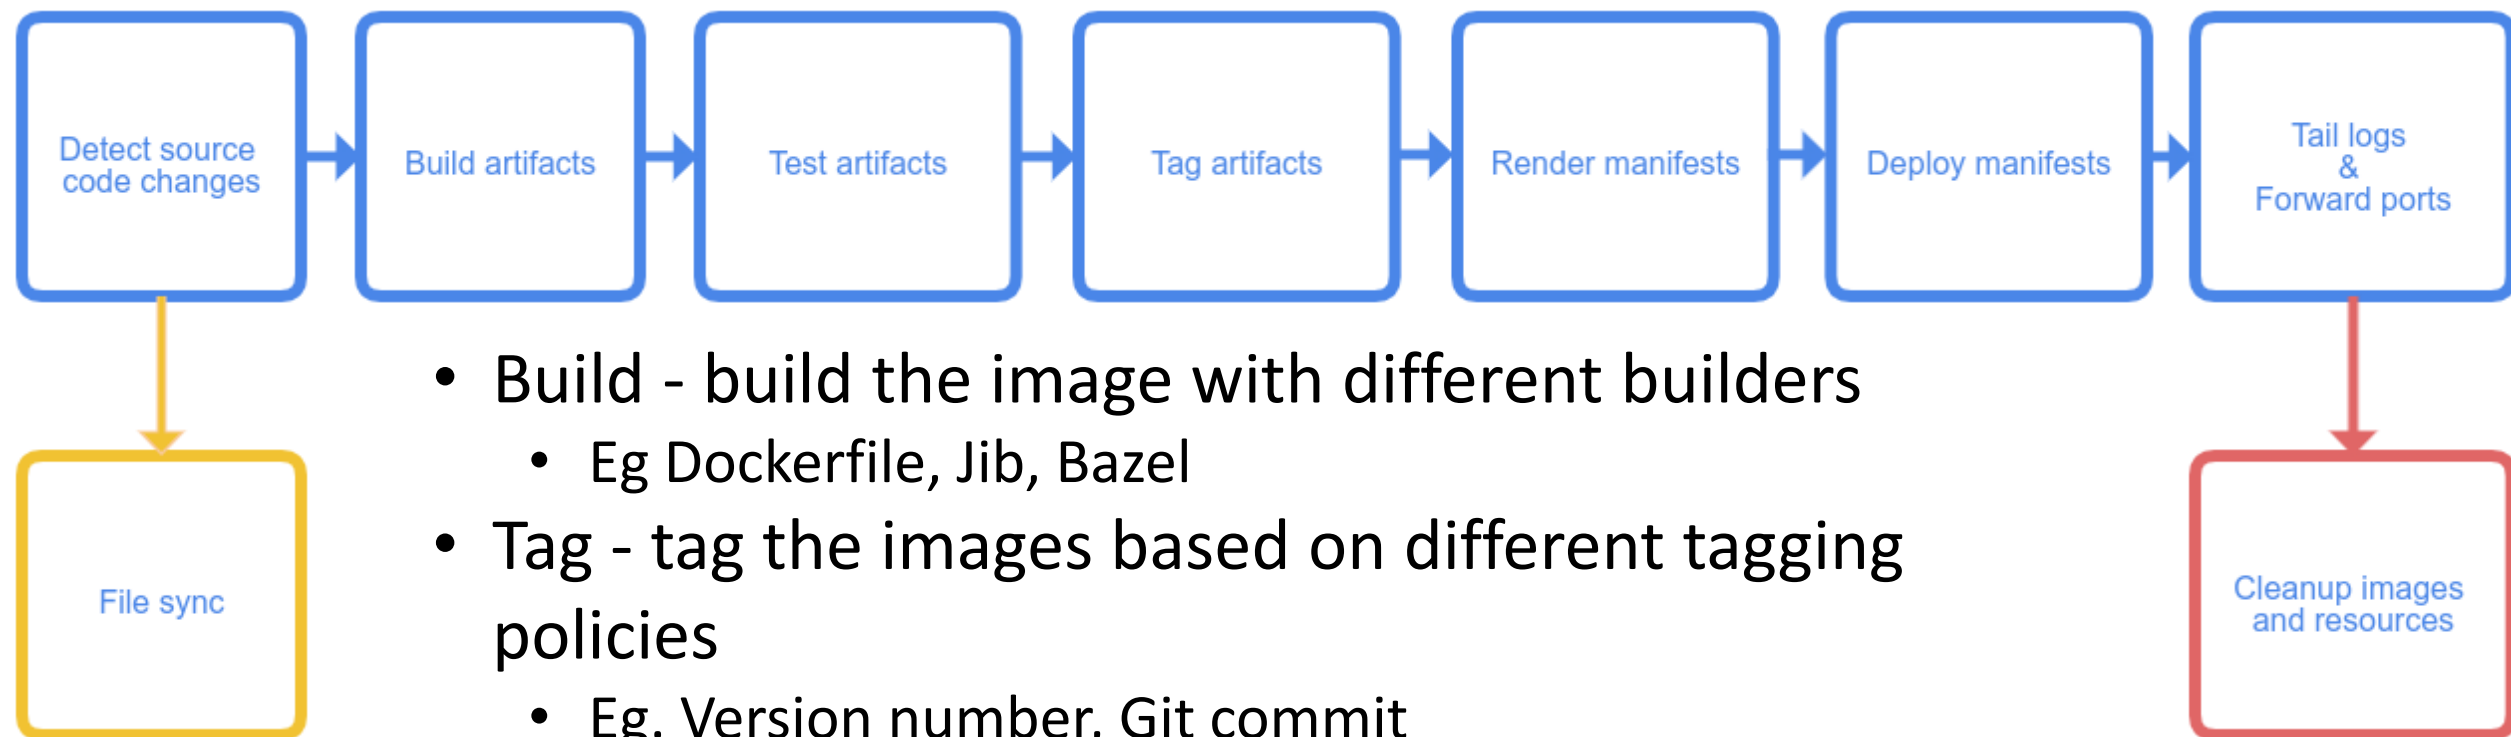


Skaifold

- Skaifold is a tool to automate developing, testing and deploying applications to Kubernetes
- Watches for file changes in applications, Dockerfile and Kubernetes manifests
- Automate the development process
 - Build, Deploy, Access
 - Can use Skaifold to just execute one of the stage eg. Build
- skaffold CLI is used to automate the flow
 - Reads a default file called `skaffold.yaml`
 - Download from <https://skaffold.dev/docs/install/>



Skafold Pipeline



- Build - build the image with different builders
 - Eg Dockerfile, Jib, Bazel
- Tag - tag the images based on different tagging policies
 - Eg. Version number, Git commit
- Test - run tests
- Deploy - deploy the application
 - With kubectl, Helm or Kustomize
- Tail logs and forward ports



Skaffold Manifest

apiVersion: **skaffold/v2beta29**

Version from
<https://skaffold.dev/docs/references/yaml/>

kind: **Config**

Config kind

metadata:

name: **myapp**

Config name

build:

Build specification

deploy:

Deployment strategy

test:

Run test

portForward:

Ports to expose



Building Images

- Supports multiple tools for building images
 - Dockerfile
 - Jib (Java image builder) for Maven and Gradle
 - Buildpacks - see <https://buildpacks.io>
 - Custom script
- Docker build
 - Can build images locally or in-cluster
 - Images can be optionally pushed to registry
 - Tagging policy based on SHA, Git commit, version,



Example - build

```
build:
  artifacts:
    - image: myapp
      context: .
      docker:
        dockerfile: Dockerfile

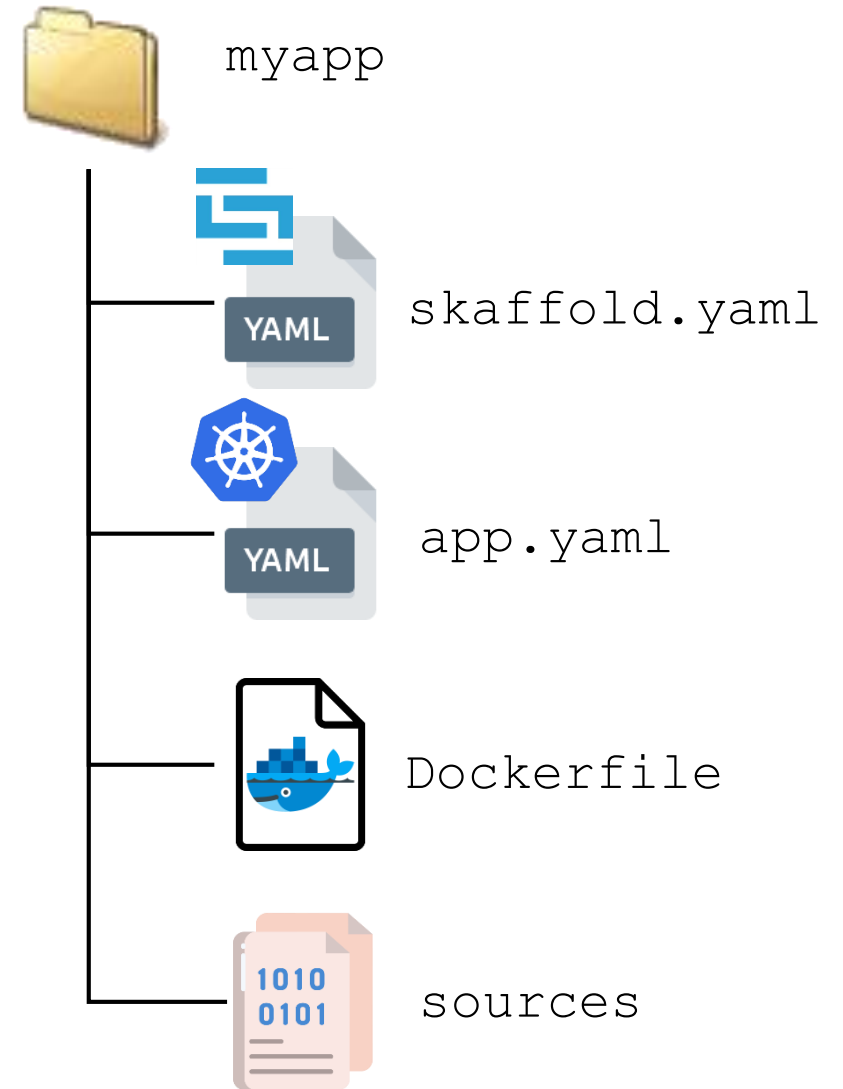
local:
  tryImportMissing: true
  tagPolicy:
    envTemplate:
      template: "{{ .APP_VERSION }}"
```

Import images that are not present

Tag the image with the environment variable

With skaffold

```
APP_VERSION="v1" skaffold build
```



With Docker

```
docker build -t fred/myapp:v1
docker push fred/myapp:v1
```



Deployment

- Supports the following ways to deploy the application
 - kubectl
 - helm
 - kustomize
- CLI for all of the above must be installed on the machine that runs the scaffold command



Example - deploy with kubectl

```
deploy:
  kubectl:
    manifests:
      - app.yaml
  defaultNamespace: myapp-ns
```

One or more Kubernetes resource files

manifests supports patterns for. Example:
all the files with k8s prefix and .yaml suffix

`k8s-*.yaml`

all the files in the k8s directory

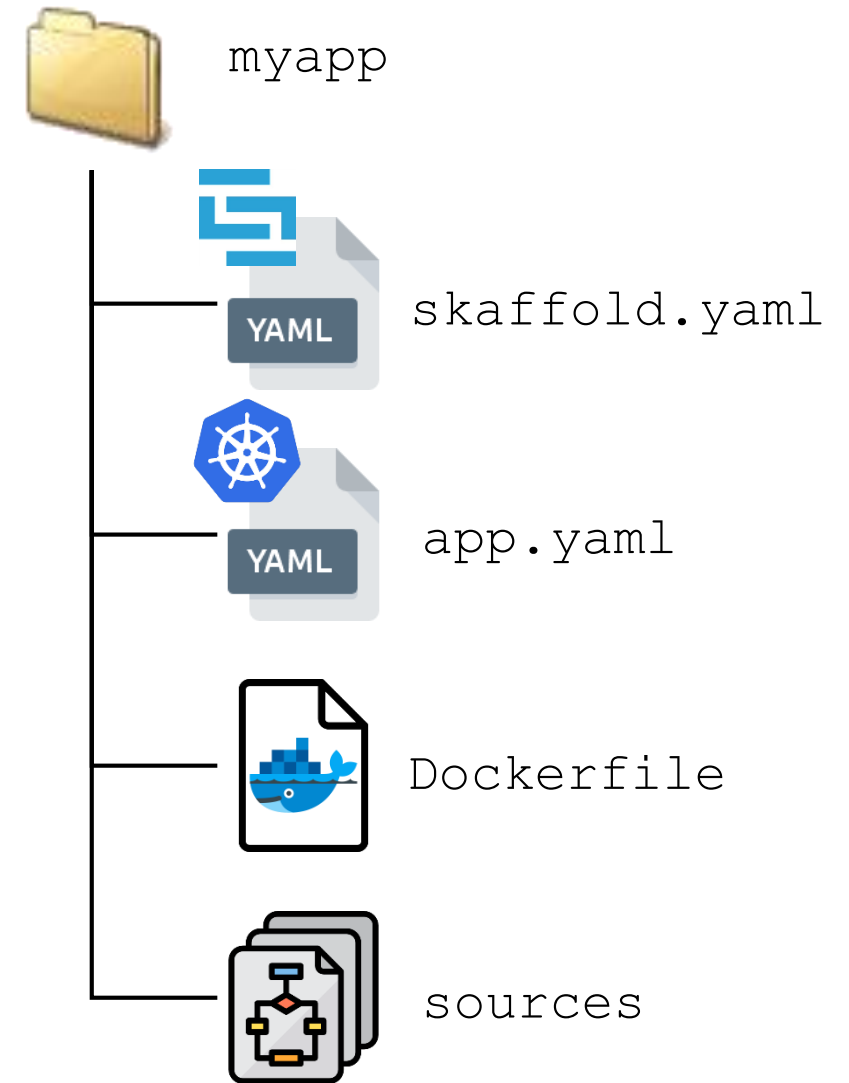
`k8s/*.yaml`

With skaffold

```
APP_VERSION="v1" skaffold run
```

With kubectl

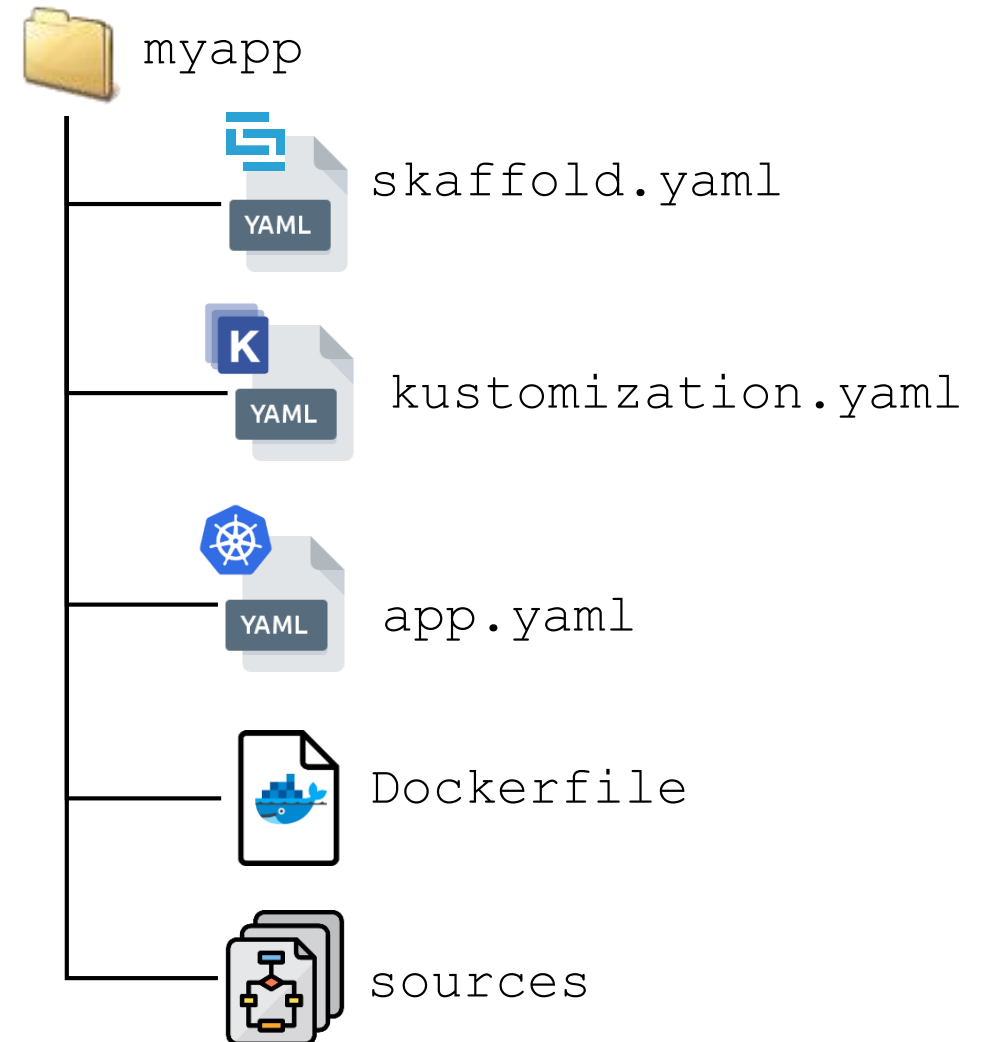
```
kubectl apply -f app.yaml
```





Example - deploy with Kustomize

```
deploy:
  kustomize:
    paths:
      - ./kustomization
    defaultNamespace: myapp-ns
```





Example - Port Forward

portForward:

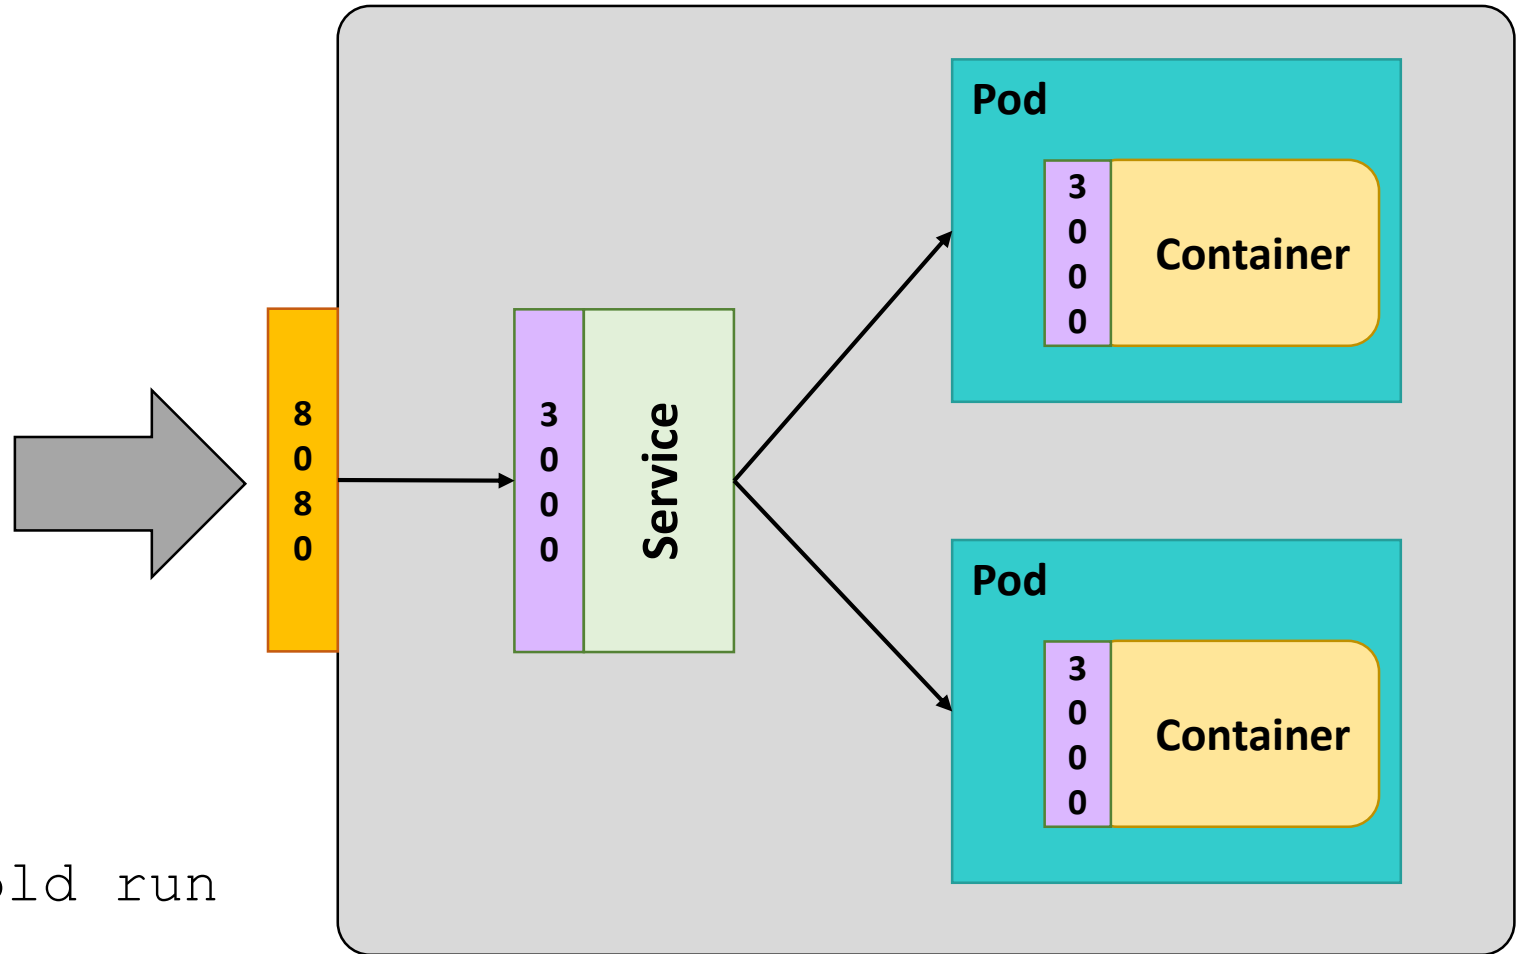
- resourceType: Service
- resourceName: app-svc
- namespace: myapp-ns
- port: 3000
- localPort: 8080

With skaffold

```
APP_VERSION="v1" skaffold run
```

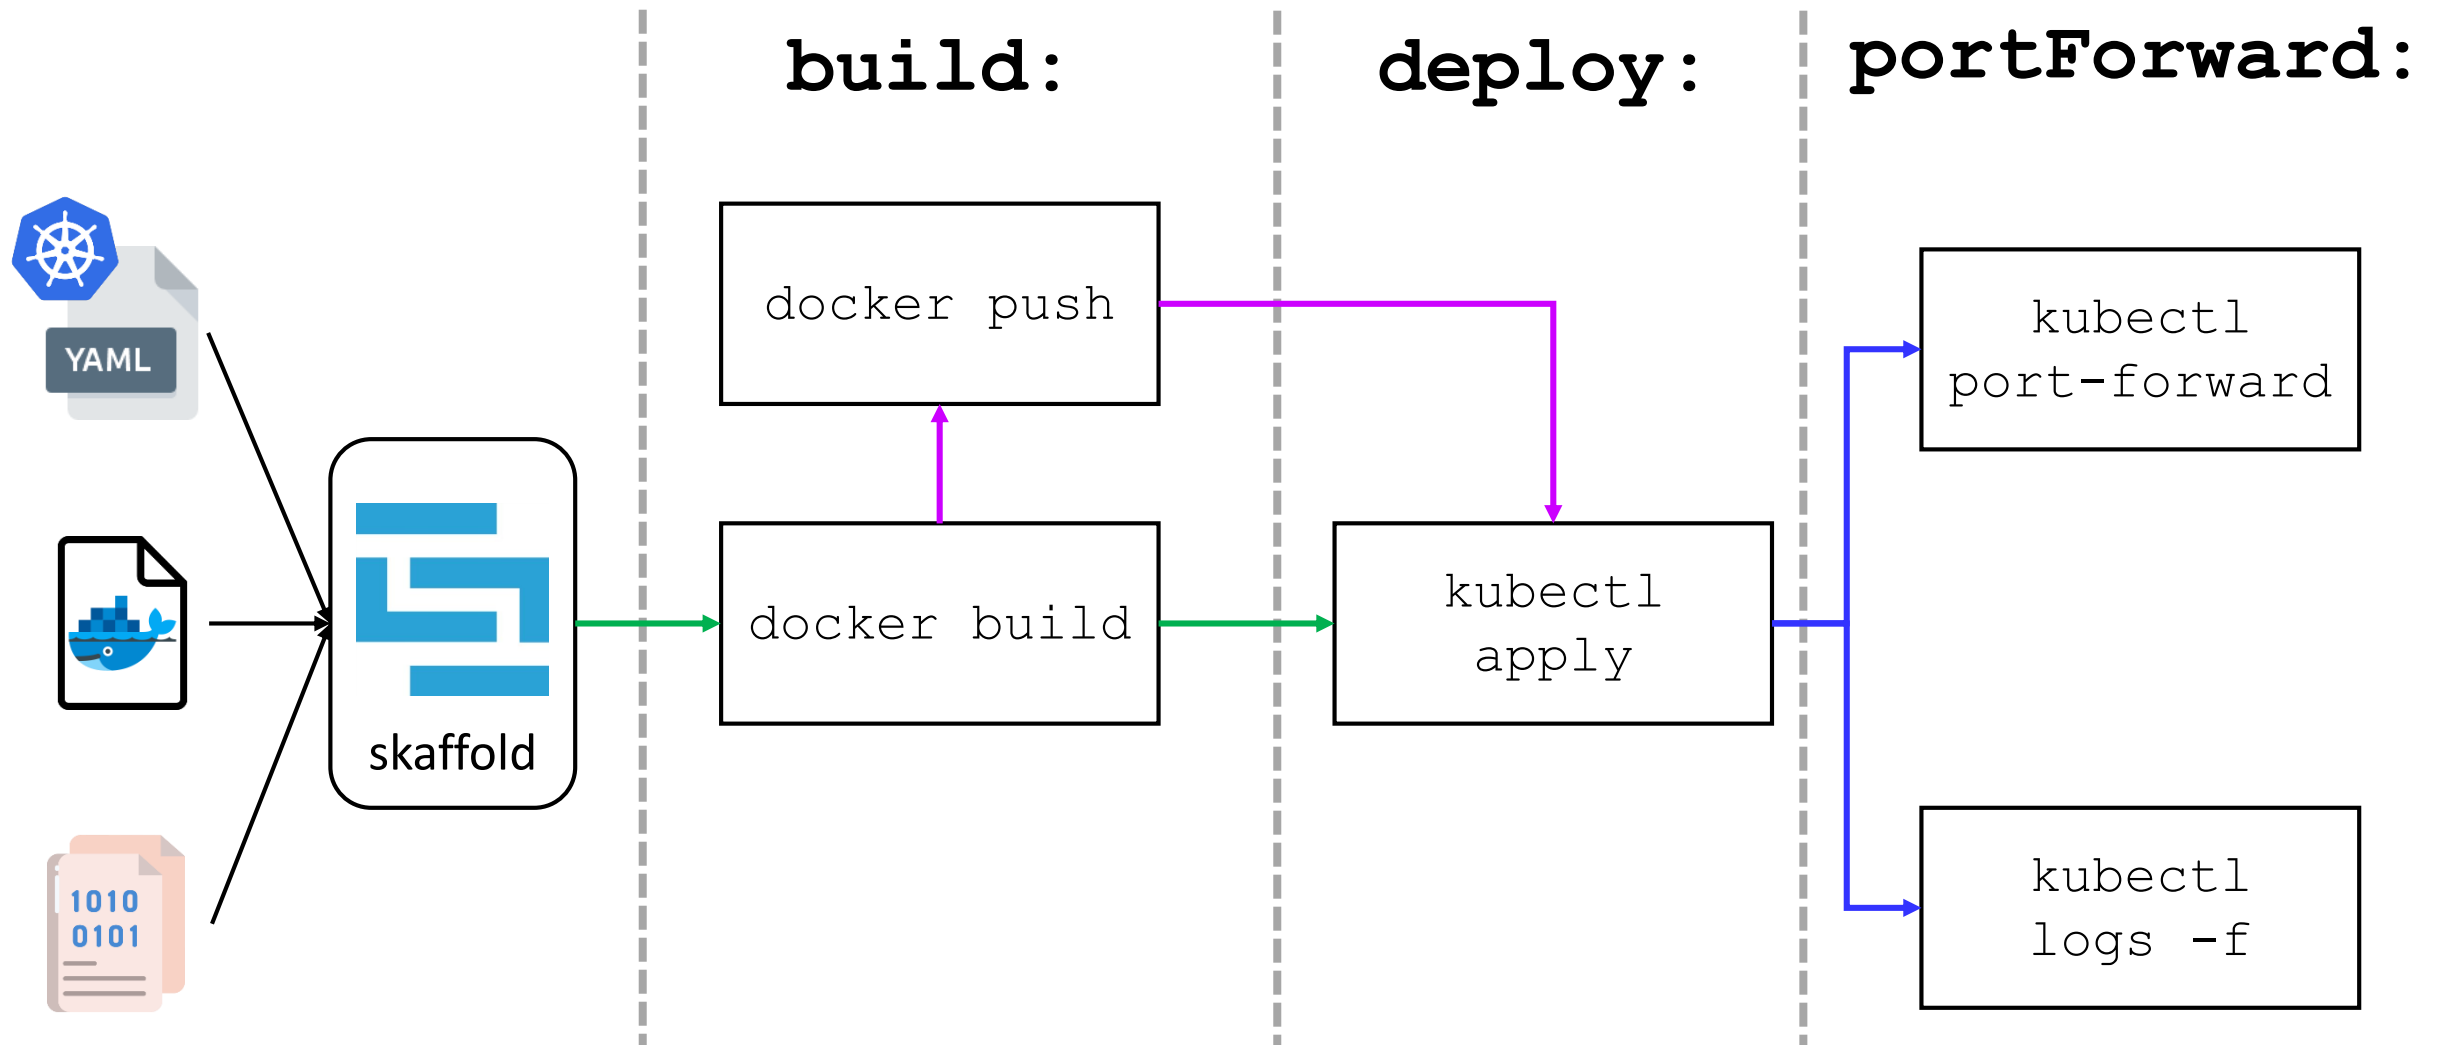
With kubectl

```
kubectl port-forward svc/myapp-svc 8080:3000 -n myapp-ns
```






Skaffold Pipeline






Skaffold Commands


`skaffold build`



`skaffold run`



`skaffold dev`





Working with Monorepos

Each service has its own
Dockerfile, Kubernetes
manifests and `scaffolding.yaml`

`scaffolding.yaml`

```
apiVersion: scaffolding/v2beta29
kind: Config
metadata:
```

```
  name: m-services
```

```
requires:
```

- path: ./m-svc0
- path: ./m-svc1

Use the `scaffolding.yaml` file
in the respective directories

