



Kubernetes

Part 3



Network Admission

- Refers to the traffic between
 - Pods in the cluster
 - External into the cluster
- Approaches to traffic control
 - Blacklist - default is to allow access unless its prohibited
 - Whitelist - default is to block access unless its permitted
- Possible approach to take for
 - Internal - trusted environment, blacklist
 - External - untrusted environment, whitelist



In Cluster Traffic

- Pods are not isolated when they are deployed
 - Default behaviour
 - Accept traffic from any sources - other pods and external entities
 - Connect to any destinations - other pods and external entities
- May want to restrict ingress and egress traffic
 - Eg. multi tenancy - only allow traffic for pods within a tenant but not to other tenants and visa versa
 - Eg. micro services - only allow pods that owns the database to connect to it
- Kubernetes uses network policy resource to manage traffic
 - Need to install a network plugin that supports network policy



Network Policy

- Network policy define a set of rules to regulate how traffic flow between a pod and its environment
- Network policy only has 'allow' rules, no deny rule
- Once a policy has been applied to a pod/set of pods
 - All traffic is blocked unless there is an allow rule
 - Eg cannot receive traffic from Ingress controller unless there is a rule that allows it - ingress rule
 - Eg. cannot query DNS unless there is a rule that allows it - egress rule
- Rules are additive
- Evaluated by ORing the rules, not AND
 - Evaluation are not affected by the order of the rules
 - Most firewalls process rules in the order which they are defined eg. iptables, ufw



Basic Network Policy

Network policy is a scoped resource. The rule applies to pod(s) in the specified namespace

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: app-netpol
  namespace: myns
spec:
```

podSelector:

```
  matchLabels:
    name: webapp
```

Select all pods in the namespace that matches this selector

policyTypes:

- Ingress
- Egress

ingress:

...

egress:

...

Ingress and Egress rules

Policy types. Can be Ingress, Egress or both



Rules

- Every ingress/egress rules has the following attributes
 - `from` - for Ingress
 - `to` - egress
 - `ports` - list of allowable ports to connect from or to
 - Allow connection from/to ports if ports is not specified



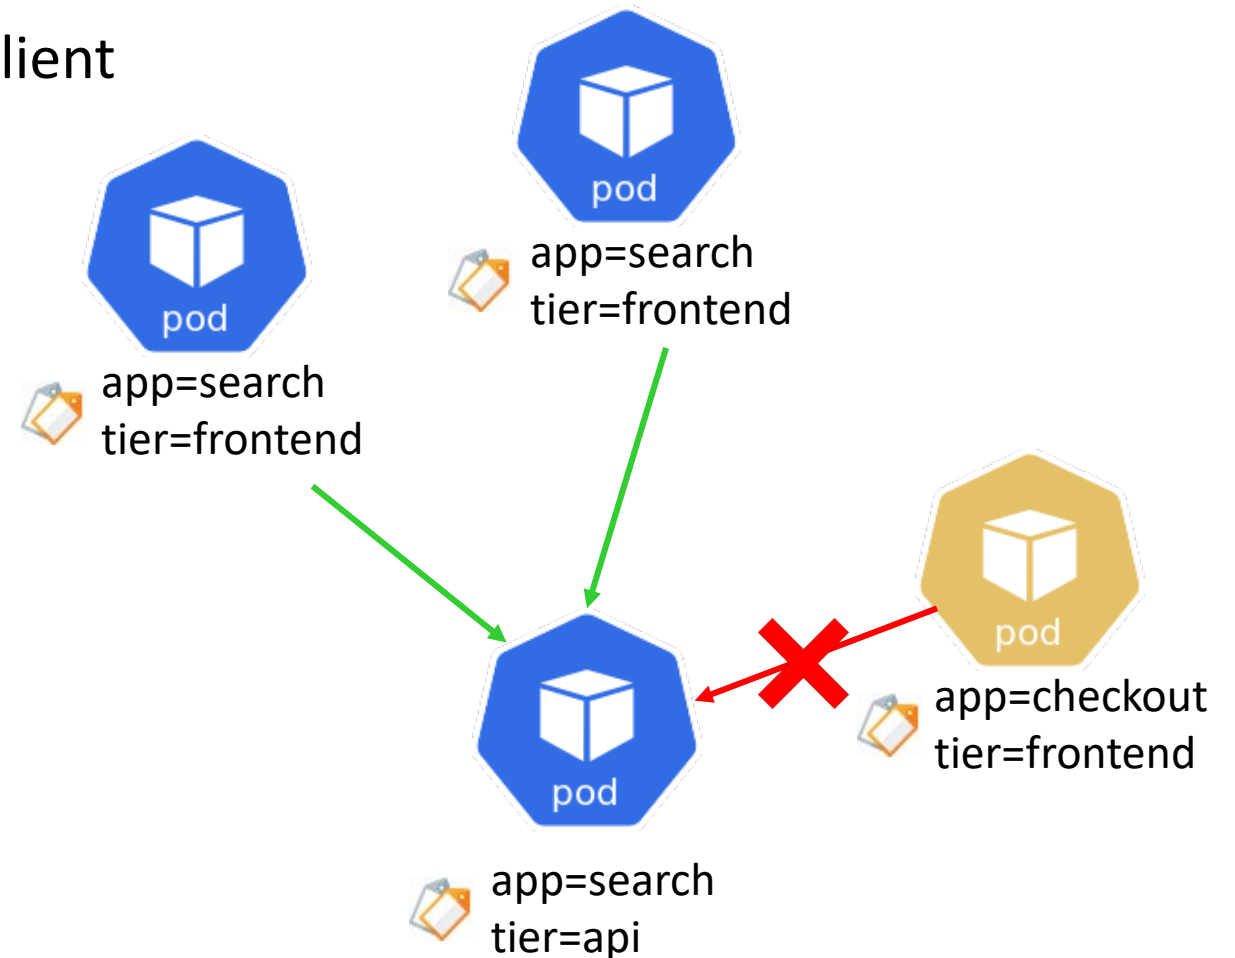
- Ingress/egress selectors
 - `podSelector`
 - Select pod(s) in the local namespaces
 - `namespaceSelector`
 - Select all pods in a given namespace
 - `namespaceSelector` and `podSelector`
 - Select specific pods from a given namespace
 - `ipBlock`
 - Select ingress/egress based on IP CIDR ranges



Network Policy Example 1

Allow only traffic from an application's client

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  ingress:
    from:
    - podSelector:
        matchLabels:
          app: search
          tier: frontend
```



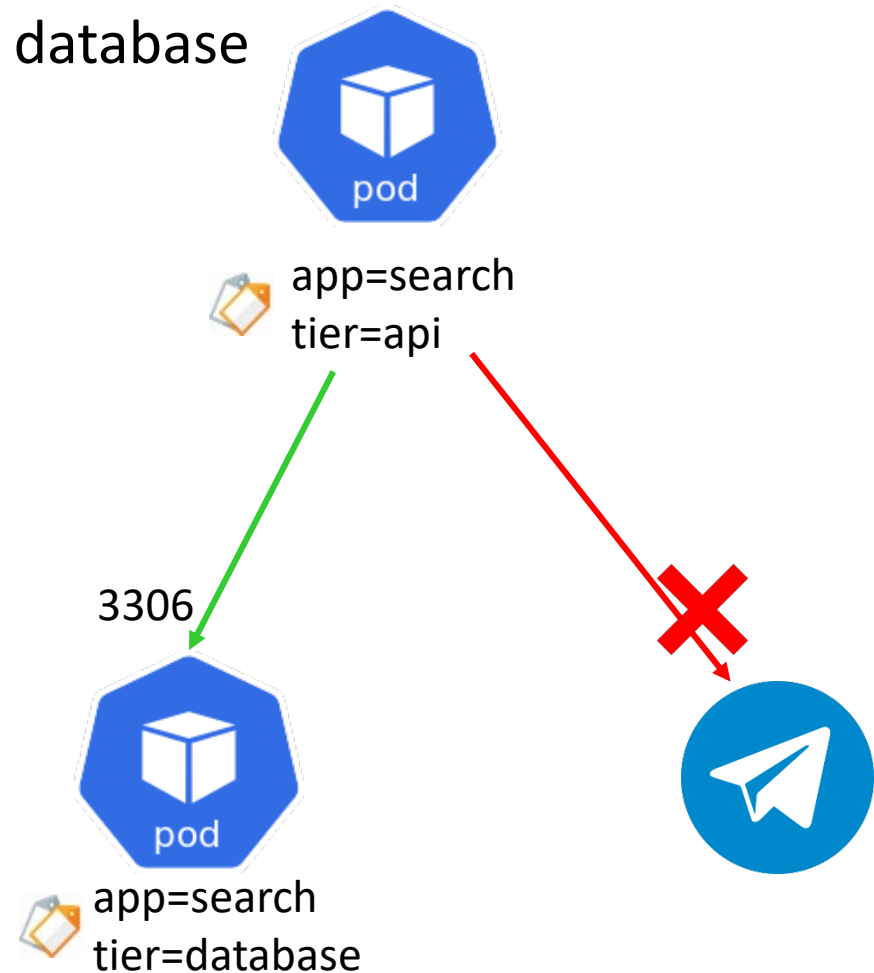
Assume pods are in `mys`



Network Policy Example 2

Allow outbound traffic only to the application's database

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Egress
  egress:
    to:
    - podSelector:
        matchLabels:
          app: search
          tier: database
  ports:
  - port: 3306
    protocol: TCP
```



Assume pods are in `myns`

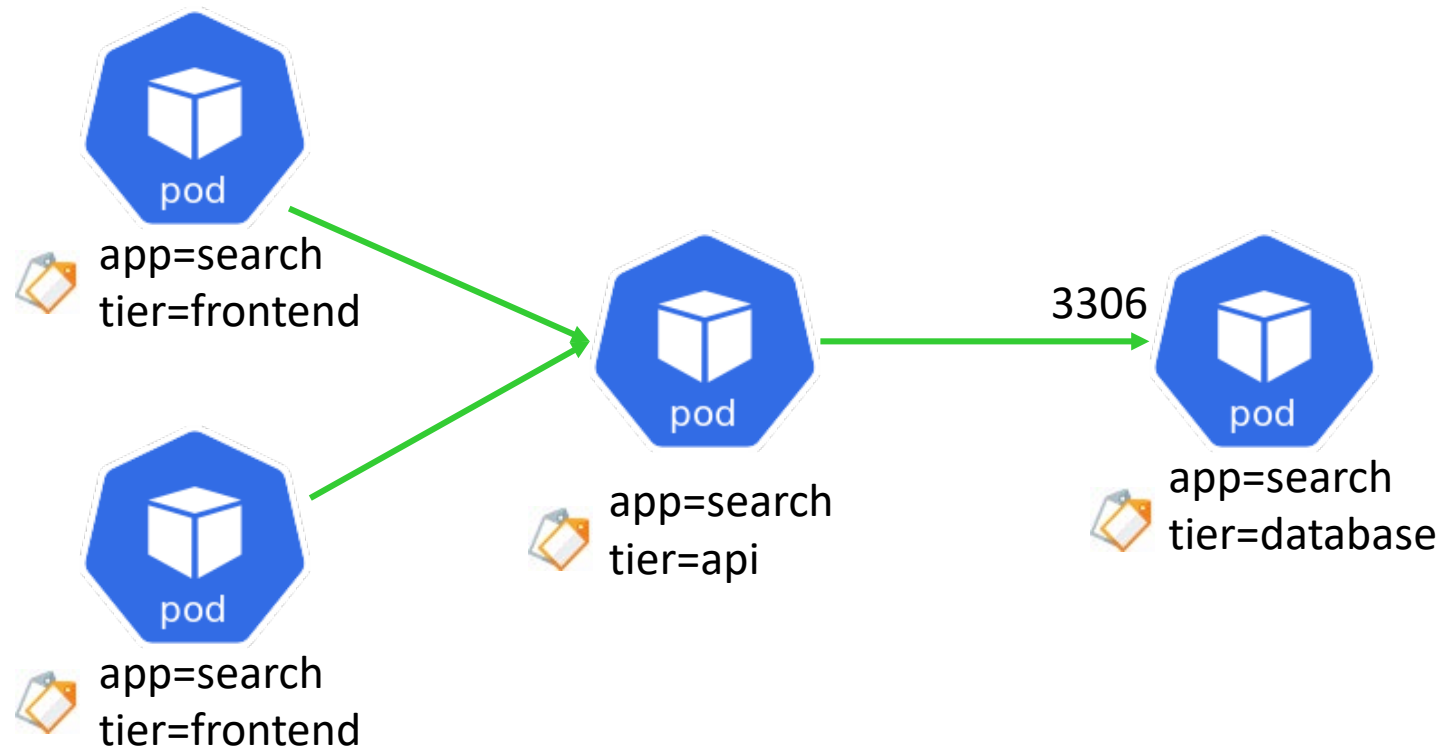


Network Policy Example 3

spec:

```
podSelector:
  matchLabels:
    app: search
    tier: api
policyTypes:
- Ingress
- Egress
ingress:
  from:
  - podSelector:
      matchLabels:
        app: search
        tier: frontend
egress:
  to:
  - podSelector:
      matchLabels:
        app: search
        tier: database
ports:
- port: 3306
  protocol: TCP
```

Allow traffic flow between all the pods in a micro service



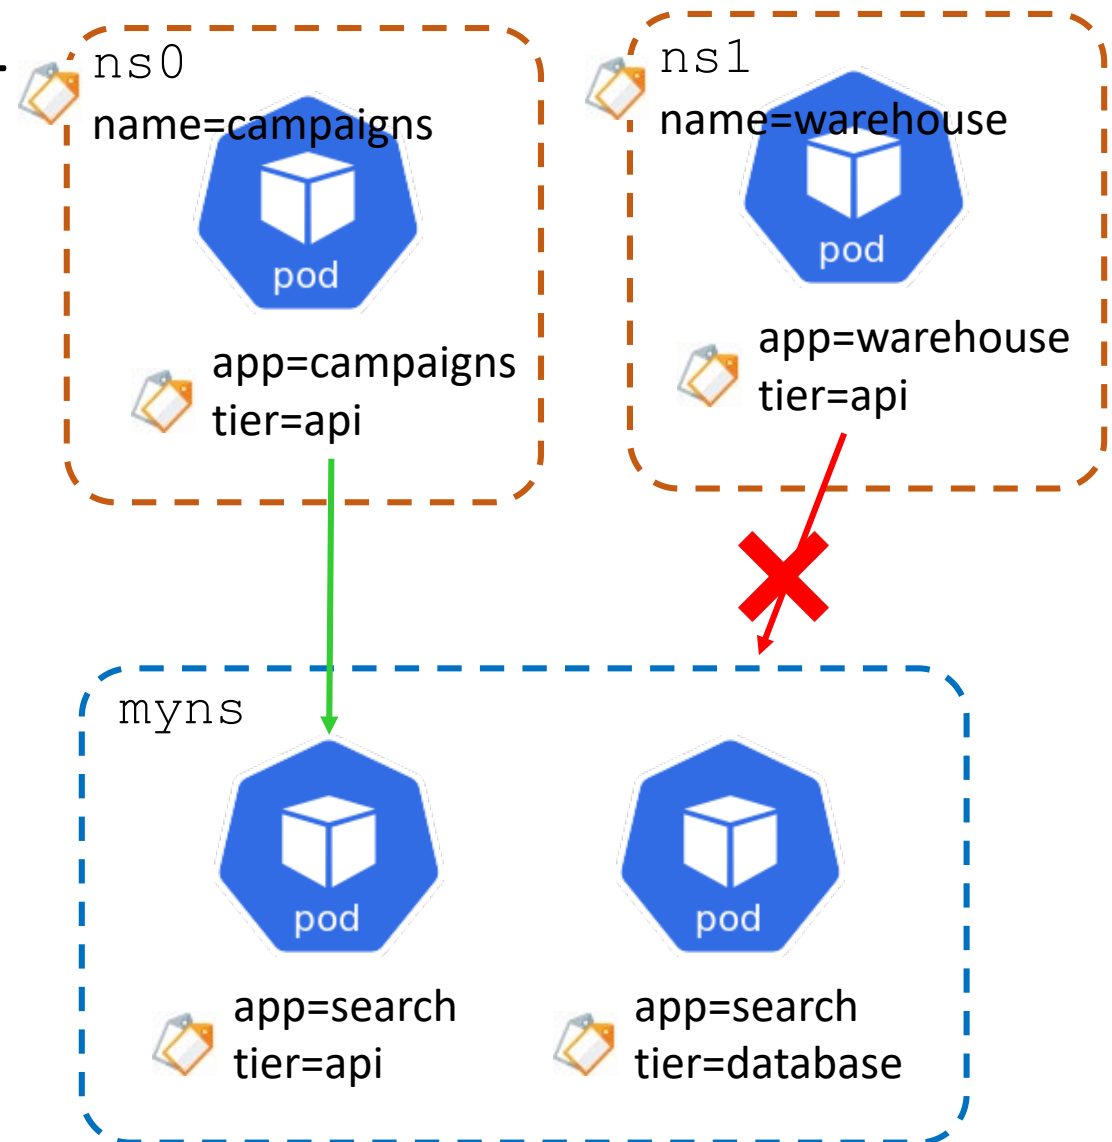
Assume pods are in `mys`



Network Policy Example 4

Allow traffic from pods of in all namespaces

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  ingress:
    from:
      - namespaceSelector:
          matchLabels:
            name: campaigns
```



Assume pods are in myns

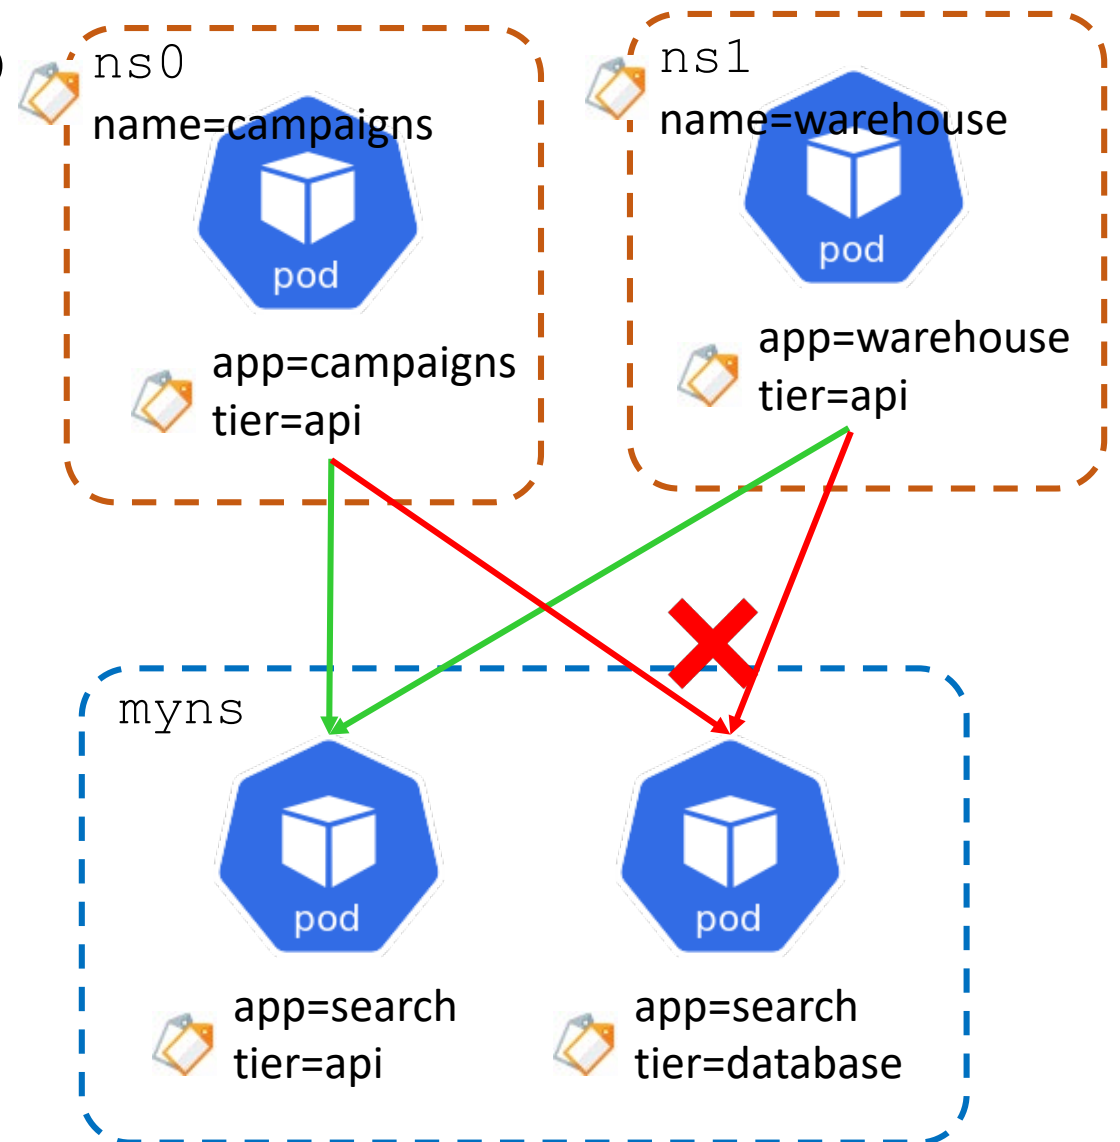


Network Policy Example 5

Allow traffic from pods of in all namespaces

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  ingress:
    from:
      - namespaceSelector: {}
```

Empty object matches any namespace



Assume pods are in myns




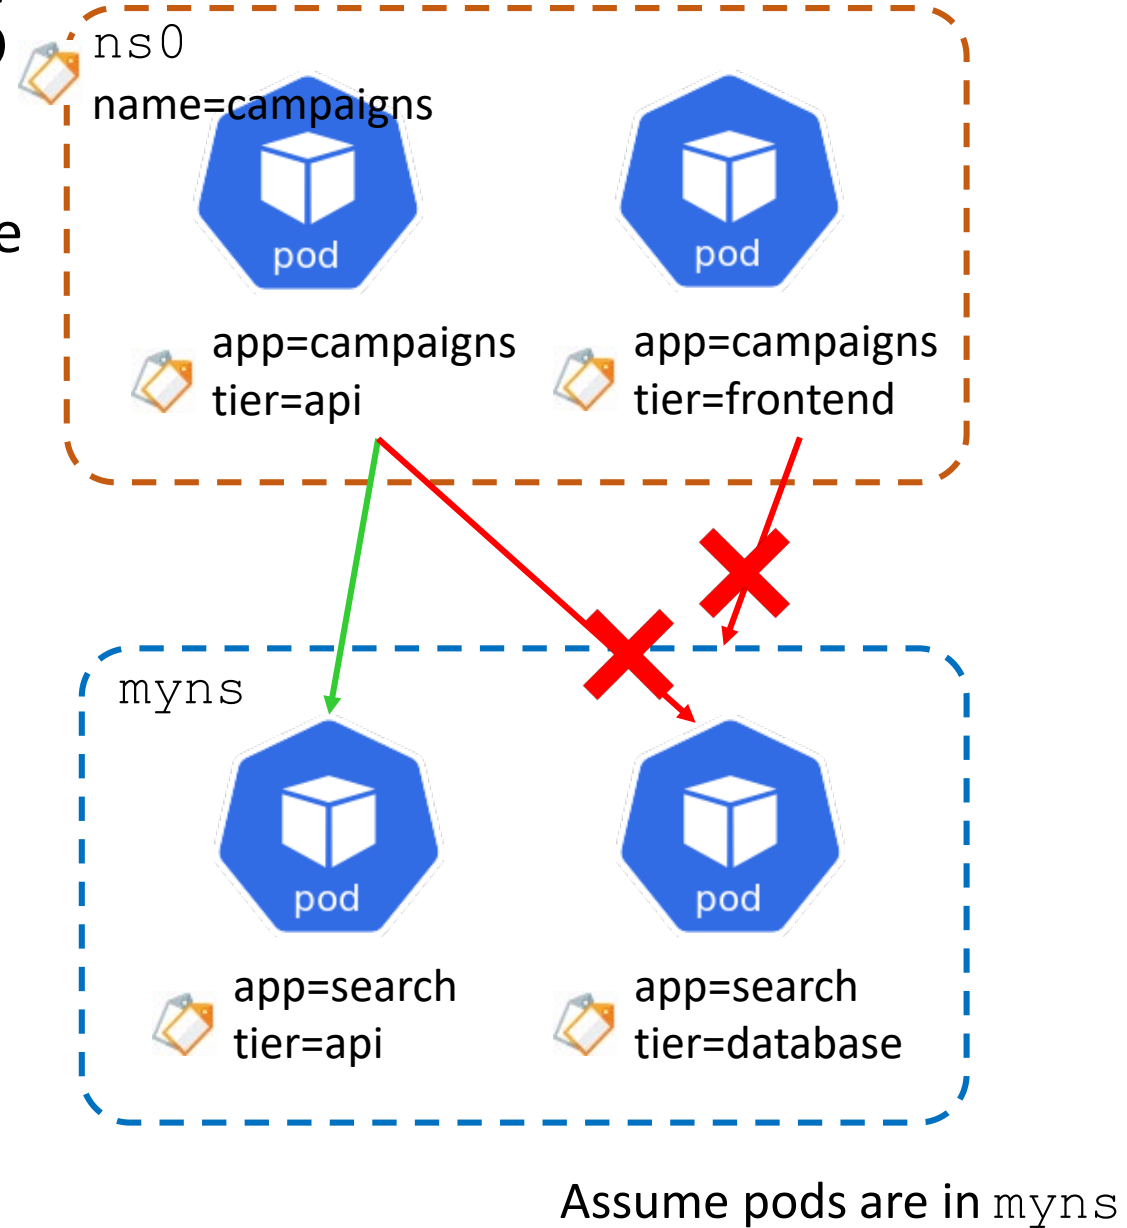
Network Policy Example 6

Allow traffic from certain pods of a namespace

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  ingress:
    from:
      - namespaceSelector:
          matchLabels:
            name: campaigns
        podSelector:
          matchLabels:
            tier: api
```

1 rule
2 conditions

 AND





Network Policy Example 7

Allow all traffic from local and another namespace

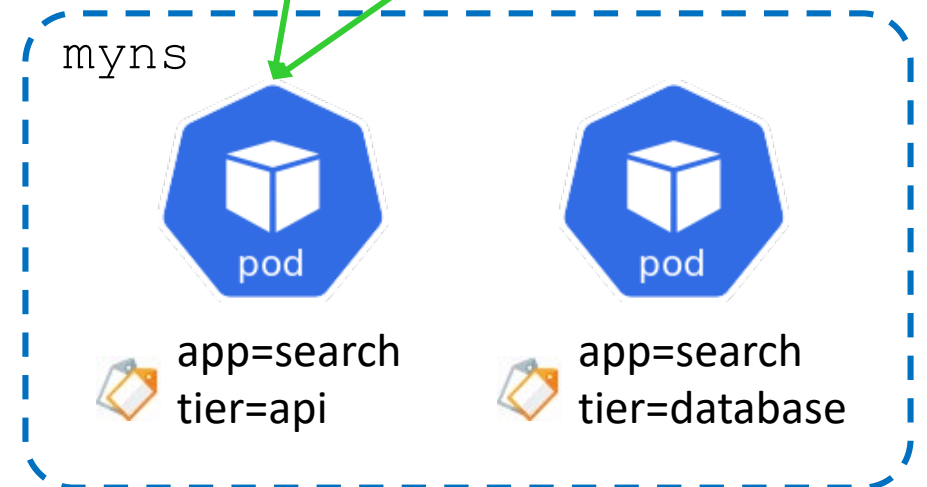
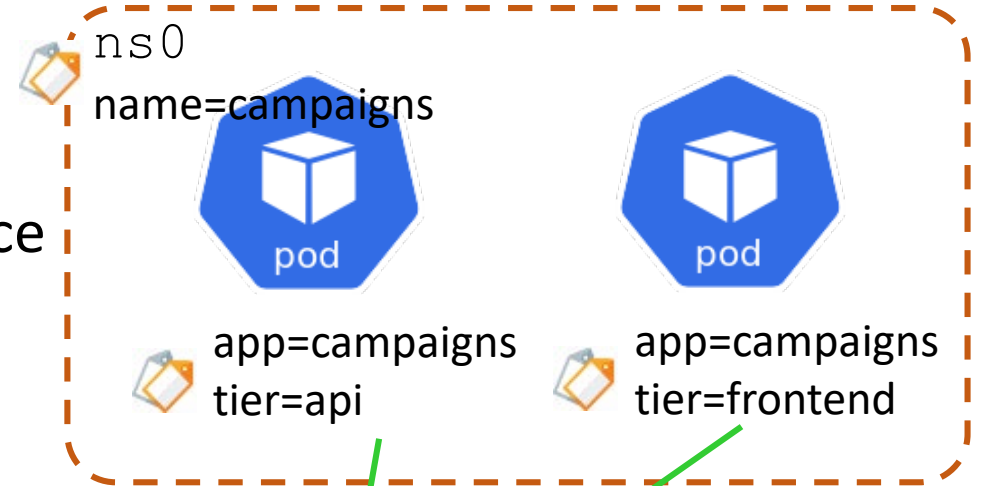
```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  ingress:
    from:
      - namespaceSelector:
          matchLabels:
            name: campaigns
      - podSelector:
          matchLabels:
            tier: api
```

2 rules
1 condition

OR

Any pods from
campaigns
namespace

Pods from local
namespace



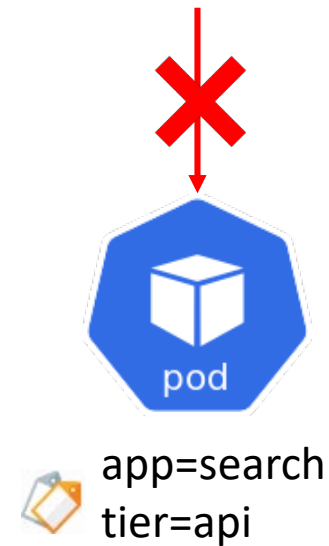
Assume pods are in myns



Network Policy Example 8

Deny all traffic

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
    - Ingress
    - Egress
```



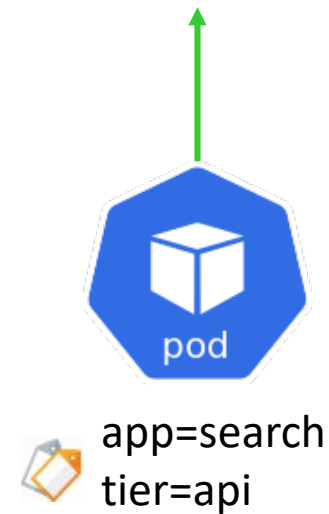
Assume pods are in `myns`



Network Policy Example 8

Allow all traffic from any source and to any destination - default behaviour

```
spec:
  podSelector:
    matchLabels:
      app: search
      tier: api
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - {}
  egress:
  - {}
```



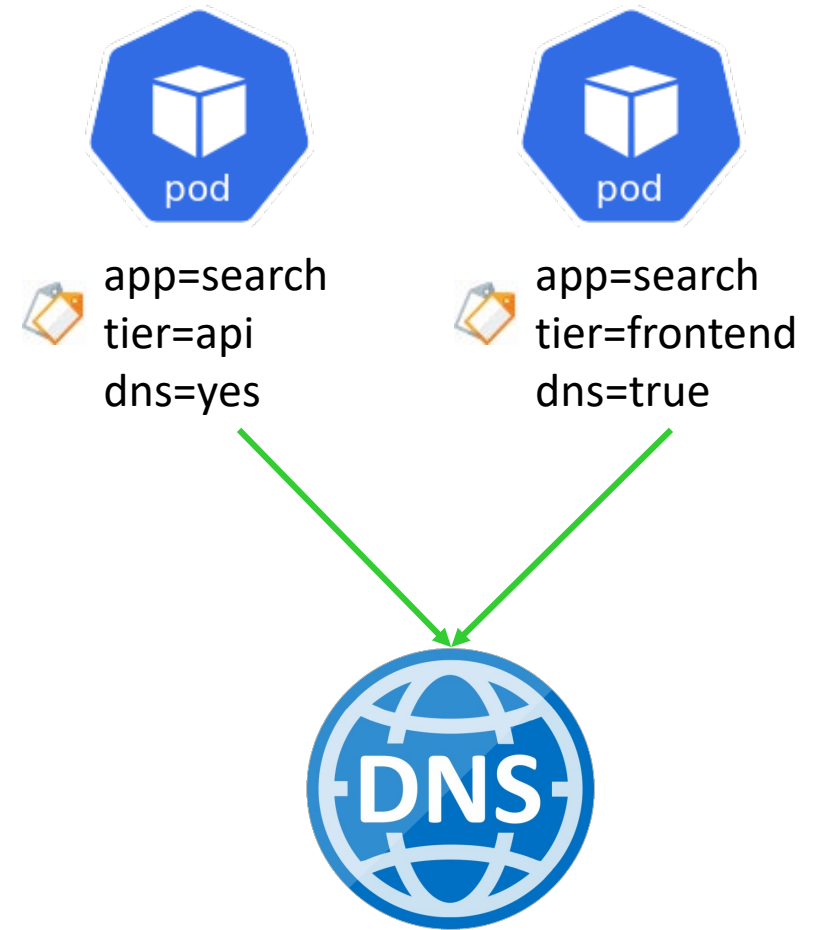
Assume pods are in `myns`



Network Policy Example 9

Allow pods to query DNS

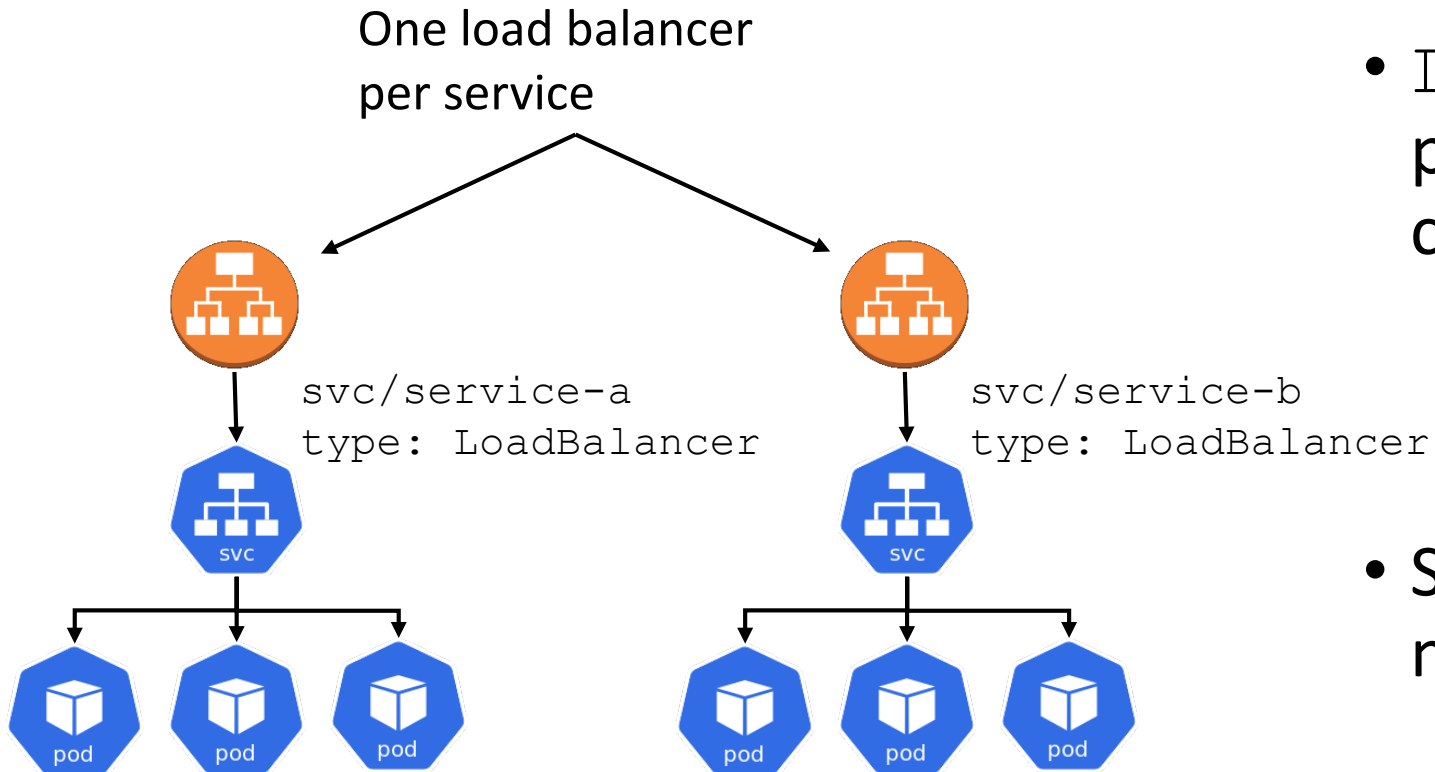
```
spec:
  podSelector:
    matchExpressions:
      - key: dns
        operator: In
        values: [ "yes", "true" ]
  policyTypes:
    - Egress
  egress:
    - ports:
        - port: 53
          protocol: UDP
        - port: 53
          protocol: TCP
```



Assume pods are in `mys`



LoadBalancer Service Type



- LoadBalancer service provision a load balancer on cloud provider
 - Allows external traffic into the cluster
 - 1 to 1 correlation
- Smaller services may not need a load balancer
 - No way of sharing
 - Not cost effective



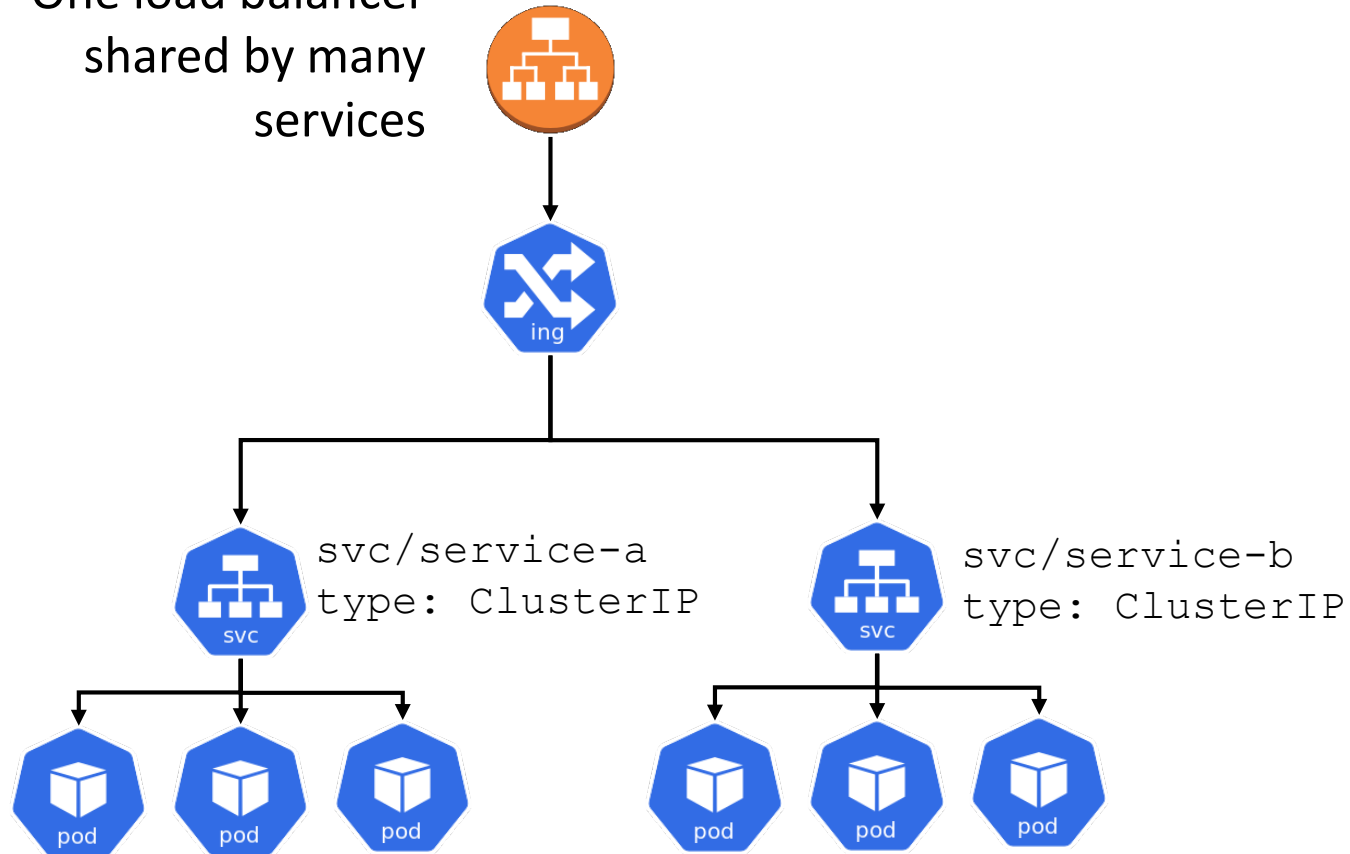
Single Entry Point

- Not secure if there are multiple uncontrolled entry points into the cluster
 - This is the case if we allow LoadBalancer service type
- Should have a single entry point
 - A choke point where all traffic enters
- Implement and enforce enterprise level policies at a single point
 - Eg. security, logging, etc
 - Certificates
- Reduce attack surface



Ingress

One load balancer
shared by many
services

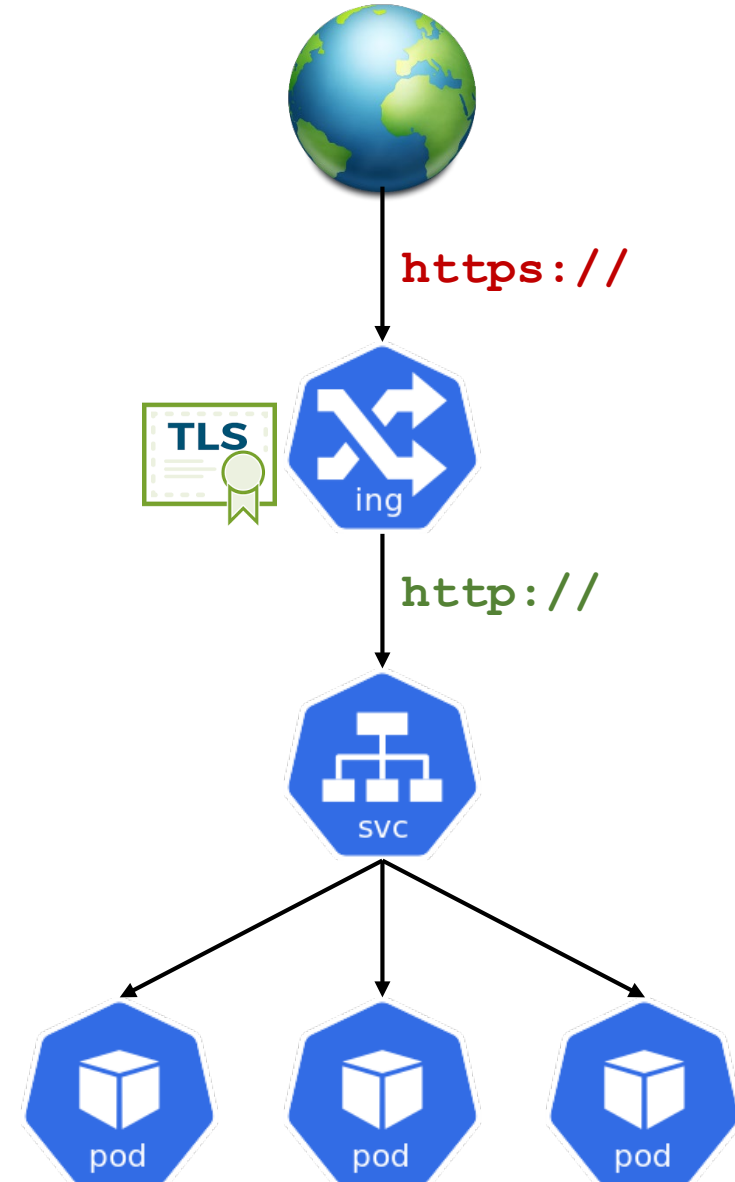


- Like LoadBalancer service type but is shared between many ClusterIP type service
 - Regulate external traffic to many different services
 - Traffic can be routed based on FQDN, resource or both
- Different ingress controllers have additional capabilities eg. OpenAPI, authentication, tracing, etc.
 - Implementation dependent



TLS

- Securing and endpoint with TLS provides the following benefits
 - Encrypt the communication channel between the client and the application
 - Protect sensitive data from being snoop or sniffed
 - Ensure that the client is talking to the 'correct' application
 - Can verify the server from its certificate
- TLS should be enabled at ingresses and not at individual application/pod
 - Better certificate and key management
- TLS terminates at the ingress
 - Traffic between the ingress and the pod is not encrypted





Enable TLS on Ingress

Example only. Generate a X.509
certificate for app.acme.com domain

↙

```
openssl req -x509 \  
-key app.key \  
-out app.cert \  
-days 1000  
-subj "/CN=app.acme.com"
```

```
kubectl create secret tls app-tls \  
--key=app.key \  
--cert=app.cert \  
-n myns
```

↖

Create a TLS secret to be used
by the ingress resource

```
apiVersion: networking.k8s.io/v1  
kind: Ingress  
...  
spec:  
  ingressClassName: nginx  
  rules:  
    - host: app.acme.com  
      http:  
        paths:  
          ...  
      tls:  
        - app.acme.com  
          secretName: app-tls
```

↗

Certificate to use for
the host. Domain name
must match the name
in the certificate